


Stage 1 (supervised): which value should a variable take?



f_θ frozen

Stage 2 (cost-tuned): which variable is cheapest to branch on?

