



Deployed cascade: sound stages first, learned components only in the last stage



Inside one search node



Neither head fixes a variable: fixing is done only by propagation and LP-probing (sound), so a wrong prediction costs backtracking and nothing else.