


Conditioning on a partial assignment is the same as reducing the instance (Proposition 1)



Only errors on FORCED variables cost a backtrack. Their share rises from 6.3% at the root to 87.9% at depth 8 of a 10×25 instance, so the same network is trained on in-tree states (reduced instances) rather than roots; the labels are exact because S is enumerated.