

Amortized Deduction for Binary Linear Systems: Learning Branching Guidance from Exact Conditional Marginals

neuroMCTS project^a

^a*Korea University, Seoul, Republic of Korea*

Abstract

A binary linear system (BLS) asks for $\mathbf{x} \in \{0, 1\}^n$ satisfying $A\mathbf{x} = \mathbf{b}$ with $A \in \{0, 1\}^{m \times n}$, a formulation that arises in laser-based non-destructive testing, where a binary occupancy image is reconstructed from integer projections. On market-split-style instances the linear relaxation is deliberately uninformative, so backtracking search spends most of its effort on branching decisions that a cheap heuristic cannot make well. We present a learned branching heuristic for this setting and an exact procedure for training it.

The method rests on two observations. First, conditioning a BLS on a partial assignment is *identical* to substituting the assignment out and obtaining a smaller BLS; a single network therefore predicts at the root and at every interior node without architectural change, and training states are ordinary instances of reduced size. Second, on instances small enough to enumerate, the exact conditional marginal $p_j = \Pr[x_j = 1 \mid A, \mathbf{b}]$ is computable, and it partitions the free variables of a node into those every surviving solution agrees on — where branching the other way empties the subtree and forces a backtrack — and those solutions disagree on, where either branch still reaches a solution. Only the first kind carries cost. This partition supplies both an exact training target and a criterion for which predictions matter: at the root only 6.3% of variables are of the first kind, but by depth 8 the share is 87.9%, which is where prediction accuracy begins converting into search cost.

We train a 21.8K-parameter bipartite graph network on these targets and use it for branching only, never for fixing variables, since it produces a probability rather than a proof. Against LP-probing — the sound procedure that detects the same variables, at one linear program each — the network needs

a single relaxation and one forward pass, so their cost ratio grows linearly in instance size: $1.7\times$ at $n = 25$, $7.2\times$ at $n = 60$, $17.8\times$ at $n = 150$, measured with a warm-started backend on both sides. Under a matched wall-clock budget on a warm-started search loop, replicated over three seeds, the resulting search is $2.6\text{--}3.0\times$ faster in median than LP-guided branching at 21×60 (faster on 23/30 instances in every seed, sign test $p = 0.005$) and expands $9.5\times$ fewer nodes, with both strategies solving every instance; within the full pipeline, where lattice reduction runs first, it closes the residual $3.0\times$ faster than LP-guided branching, while random branching closes none of it. Controlling the number of solutions across sizes, we find no measurable transfer penalty: a network trained only at 10×25 is statistically indistinguishable from one trained at the evaluation size. On a 100-instance 21×60 set the supervised guidance is $2.0\text{--}2.3\times$ faster than LP-guided branching with $5.5\times$ fewer nodes. Two further experiments locate what the supervised target does and does not capture. Training on states of deployment size — the regime where the network’s prediction margin lives — does *not* help: it loses competence on the small interior states where most nodes are spent, and the search gets slower. What does help is optimizing the objective the search actually pays: with the exact subtree cost of every branching decision available from the search itself, an actor-critic fine-tuning of the branching policy, initialized at the supervised rule, expands $2.3\times$ fewer nodes and is $2.0\text{--}2.2\times$ faster than the supervised policy on the same 100 instances ($p < 0.01$), $4.8\text{--}5.9\times$ faster than LP-guided branching. The fine-tuned policy branches on a different variable than the confidence rule at 95% of held-out states, typically one it is *less* sure about, which is direct evidence that the probability of being right and the cost of being wrong are different objectives.

We also report what the method does not achieve. CP-SAT solves the same instances $9\text{--}40\times$ faster at every size tested ($9\text{--}10\times$ at 21×60 after fine-tuning), and our system does not compete with it; decomposing that gap shows our branching produces $2.2\times$ *fewer* nodes but processes them $32\times$ more slowly, which places the deficit in per-node cost rather than in branching quality, though we do not show that it is recoverable. And the cost advantage over probing, while asymptotically real, is only $1.3\text{--}1.7\times$ at the depths where prediction matters most on our instances, where probing additionally offers proofs rather than predictions. The contribution we claim is therefore methodological: exact conditional posteriors are an obtainable and well-posed supervision signal for branching, in-tree states are a better training distribution than root states — a root-trained network is worse than

relaxation rounding when deployed in the tree, an in-tree-trained one is better by 3.4 points — and the resulting heuristic transfers across instance size. Code, generators, exact-marginal labels and per-instance results are released.

Keywords: Binary linear systems, Learning to branch, Amortized optimization, Exact conditional marginals, Market split, Graph neural networks

1. Introduction

1.1. Problem

Let $A \in \{0, 1\}^{m \times n}$ and $\mathbf{b} \in \mathbb{Z}^m$. A *binary linear system* (BLS) is the feasibility problem

$$\text{find } \mathbf{x} \in \{0, 1\}^n \text{ such that } A\mathbf{x} = \mathbf{b}, \quad (1)$$

whose solution set we write $\mathcal{S} = \{\mathbf{x} \in \{0, 1\}^n : A\mathbf{x} = \mathbf{b}\}$. Each row is a cardinality constraint $\sum_j a_{ij}x_j = b_i$ stating how many cells of a designated subset are occupied.

The motivating application is laser-based non-destructive testing. A specimen is scanned along m directions; each scan returns the integer count of absorbing cells along its path, and the task is to reconstruct the binary occupancy image. Two properties of that setting shape the problem formulation. Solutions need not be unique — several images can be consistent with the same measurements, and recovering *any* member of $\mathcal{S}$ solves the physical problem. And measurement noise can perturb $\mathbf{b}$ so that $\mathcal{S} = \emptyset$, making a reliable infeasibility verdict as valuable as a solution.

Deciding $\mathcal{S} \neq \emptyset$ is NP-complete in general. The instances studied here are drawn from a deliberately hard regime modelled on market-split problems [6, 2], in which the linear relaxation

$$\tilde{\mathcal{S}} = \{\mathbf{x} \in [0, 1]^n : A\mathbf{x} = \mathbf{b}\} \quad (2)$$

is a large polytope with predominantly fractional vertices. Relaxation-based bounding is therefore weak, and a complete solver must decide where to branch with little guidance from the relaxation — exactly the decision this paper learns.

1.2. Where a learned component attaches, and what to train it on

A complete solver for (1) alternates deduction (constraint propagation, relaxation-based pruning) with a *branching decision*: which free variable to split on, and which value to try first. Deduction is sound and must remain so; branching is a heuristic choice, where an error costs search time but never correctness. This asymmetry makes branching the natural place to attach a learned component, and it is where learning has been most successful in mixed-integer programming [9, 7].

The difficulty is supervision. Two standard answers are available and both are unsatisfactory here. Training against a planted solution $\mathbf{x}^*$ treats one member of $\mathcal{S}$ as ground truth, which is ill-posed when $|\mathcal{S}| > 1$: nothing in $(A, \mathbf{b})$ distinguishes the planted solution from any other, so the label is a sample from the posterior rather than the posterior. Imitating an expensive expert, as in strong-branching imitation, requires an expert whose decisions are both informative and affordable to collect; the natural candidate here, LP-probing, costs one linear program per variable per node.

We take a third route, available because these instances have planted solutions and because feasibility can be enumerated at moderate size: compute the *exact conditional posterior* by exhaustive enumeration on small instances, and train against it. Section 3 shows that this is both tractable and, in a precise sense, the right target — and that the posterior additionally identifies *which* predictions can affect search at all.

1.3. Contributions

1. **A reduction identity that makes in-tree supervision trivial to implement.** Conditioning (1) on a partial assignment yields a smaller BLS with the same structure (Proposition 1). One architecture, one label generator and one training loop therefore serve the root and every interior node, and a depth- d training state is stored as an ordinary instance on $n - d$ variables.
2. **A partition of branching decisions by whether they can cost anything.** The exact marginals split a node’s free variables into DET (all surviving solutions agree; branching the other way empties the subtree) and OPEN (solutions disagree; either branch reaches a solution). Only DET errors cost a backtrack. We measure the composition: 6.3% DET at the root, rising to 87.9% at depth 8 (Sections 4.3, 4.4). This locates useful prediction inside the tree rather than at the root and explains why root-level accuracy is a poor objective.

- 60 3. **A cost characterization of the learned heuristic against the**
61 **sound procedure it approximates.** The network needs one relax-
62 ation and one forward pass; LP-probing needs one relaxation per vari-
63 able. Measured with a warm-started backend on both sides, the ratio
64 scales from $1.7\times$ at $n = 25$ to $17.8\times$ at $n = 150$ (Section 4.5). We
65 report this as a scaling law rather than a single figure, and note that
66 at the sizes where our search operates the advantage is modest and
67 probing has the additional merit of soundness.
- 68 4. **End-to-end evaluation under matched wall-clock budgets.** At
69 21×60 on a warm-started search loop, replicated over three seeds, the
70 learned guidance is $2.0\text{--}2.3\times$ faster in median than LP-based branching
71 on a 100-instance set (faster on 70–71/100, sign test $p < 0.001$) and
72 expands $5.5\times$ fewer nodes; on the original 30-instance set the margins
73 are $2.6\text{--}3.0\times$ and $9.5\times$, and both strategies solve every instance within
74 600 s (Section 4.6). In the deployed cascade both sound strategies close
75 every residual instance, the learned one $3.0\times$ faster (40/50 wins pooled
76 over seeds, $p = 2 \times 10^{-5}$), while random branching closes none of the
77 50 residual instances (Section 4.12).
- 78 5. **Size transfer with solution multiplicity controlled.** Holding $|\mathcal{S}|$
79 comparable across sizes, a network trained only at 10×25 matches one
80 trained at the evaluation size 21×60 (faster on 15/30 instances; total
81 time differing by 0.4%), showing the learned quantity is not size-specific
82 (Section 4.8).
- 83 6. **A negative result on aligning the training distribution with**
84 **deployment.** The network’s margin over relaxation rounding lives on
85 states with 50–60 free variables, outside its 13–25 training range. Train-
86 ing on exactly those states, with everything else fixed, raises accuracy
87 there by about one point and *doubles* the search tree, because it loses
88 competence on the small interior states where two thirds of all nodes
89 are spent (Section 4.10).
- 90 7. **Cost-aware fine-tuning of the branching policy.** The exact cost
91 of a branching decision — the number of nodes its subtree expands —
92 is available from the search itself, and on held-out states the confidence
93 rule picks the cheapest of its five candidates only 22% of the time.
94 An actor-critic fine-tuning on that cost, initialized at the supervised
95 rule with the marginal network frozen, expands $2.3\times$ fewer nodes and
96 is $2.0\text{--}2.2\times$ faster than the supervised policy on the 100-instance set
97 (70/100 and 64–67/100 wins, $p < 0.01$ in every seed), $4.8\text{--}5.9\times$ faster

98 than LP-guided branching (Section 4.11).

99 2. Related Work

100 2.1. Binary linear systems and market-split hardness

101 Problem (1) is a 0/1 integer feasibility problem and is NP-complete for
102 general A . Practical hardness, however, depends on the instance distribution.
103 Cornuéjols and Dawande [6] introduced the market-split family, in which m
104 agents must divide n items so that every agent’s budget is met exactly; with
105 $n \approx 10(m - 1)$ these instances defeat branch-and-bound at remarkably small
106 sizes. The mechanism is the integrality gap: the relaxation polytope is large
107 and its vertices are fractional, so bounding prunes little and the search must
108 branch nearly blindly. Aardal et al. [2] showed that lattice basis reduction
109 attacks the same family far more effectively than branch-and-bound, a result
110 we rely on in Section 3.2.

111 Our generator (Section 4.1) makes this mechanism explicit by maximizing
112 the spread of relaxation vertices, which enlarges $\tilde{\mathcal{S}}$ and weakens the relaxation
113 signal by construction.

114 2.2. Complete solvers

115 *Mixed-integer programming.* Branch-and-bound over the linear relaxation,
116 strengthened by presolve and cutting planes, is the standard approach. Its
117 leverage comes from pruning nodes whose relaxation is infeasible or bounded
118 away from the target, which is precisely what the market-split construction
119 removes.

120 *Constraint programming with clause learning.* CP-SAT-style solvers com-
121 bine constraint propagation with conflict-driven clause learning [12, 14] and
122 lazy clause generation [15]. Their strength on this family comes from a
123 property central to our argument: a search *manufactures* information as it
124 descends. A marginal $\Pr[x_j = 1 \mid A, \mathbf{b}]$ may be uninformative at the root
125 while $\Pr[x_j = 1 \mid A, \mathbf{b}, x_1=1, x_5=0]$ is 0 or 1. Structure invisible at the root
126 becomes deducible after commitments are made, which is why Section 3.4
127 studies the interior of the tree rather than the root.

128 *Lattice reduction..* An alternative attack embeds (1) in a lattice and searches
129 for short vectors. LLL [10] reduces a basis in polynomial time and BKZ
130 [16] generalizes it with a block size trading time against basis quality. The
131 Aardal–Hurkens–Lenstra embedding [1, 2] is constructed so that a 0/1 so-
132 lution appears directly as a short vector, making reduction a certificate-
133 producing procedure rather than a heuristic. We use it as the first stage
134 of the deployed pipeline. Lattice enumeration remains the state of the art
135 on the canonical benchmark: Wassermann [35] solves QOBLIB market-split
136 instances up to $m = 14$ on one CPU where branch-and-cut stops at $m = 7$;
137 our instances are a different regime (multiplicity matched, $m/n \approx 0.35$) and
138 our contribution is a branching methodology rather than a record there.

139 2.3. Learning for combinatorial optimization

140 Bengio et al. [4] survey this field and frame it in terms we adopt: solvers
141 rely on hand-crafted heuristics for decisions that are “otherwise too expensive
142 to compute or mathematically not well defined,” and learning is a candidate
143 for making those decisions better. The phrase *too expensive to compute*
144 presumes the decision is computable in principle and asks only about cost,
145 which is the situation we measure in Section 4.5.

146 *Marginals of the solution set as a branching signal..* The fraction of solutions
147 agreeing with a decision was a branching criterion before learning: counting-
148 based search [32] branches on the variable–value pair of maximum solu-
149 tion density computed inside individual constraints, and belief-propagation-
150 guided decimation [29] fixes the most biased variable under approximate
151 marginals recomputed after each fixing. Learned versions replace the count-
152 ing by a network: NSNet [28] trains a BP-structured graph network on exact
153 marginals of a SAT formula’s satisfying assignments, obtained by enumer-
154 ation, and rounds them into a local-search initialization; MIP-GNN [27],
155 Neural Diving [30] and predict-and-search [26] predict per-variable biases
156 from collected near-optimal solutions and use them to fix variables, warm-
157 start or restrict a search. Our target is the same as NSNet’s; what differs is
158 that the marginals are conditional, carried to every node by the reduction
159 identity, and used to guide a complete search without fixing anything, and
160 that the DET/OPEN partition and the in-tree-versus-root comparison have no
161 counterpart there.

162 *Branching imitation..* The clearest success in this area is imitation of an
163 expensive expert. Strong branching scores candidates by tentatively solving

164 both children’s relaxations; it yields small trees and is too slow to run at
 165 every node. Gasse et al. [7] train a bipartite graph network to imitate it and
 166 recover most of the benefit at a fraction of the cost, following Khalil et al.
 167 [9]. The gain is not new information — the expert already knew the answer
 168 — but cheap access to it. Our method occupies the same structural position,
 169 with two differences: the target is an exact posterior rather than an expert’s
 170 scores, and the analysis identifies which of the expert’s decisions matter.

171 *Direct satisfiability prediction..* Attempts to have a network decide satisfia-
 172 bility end-to-end have a weaker record. NeuroSAT [17] classifies satisfiability
 173 by message passing over a literal–clause graph and succeeds on small ran-
 174 dom instances, but Li et al. [11] report substantial accuracy collapse under
 175 size transfer. Chen et al. [5] prove that standard message-passing networks
 176 cannot separate certain feasible/infeasible MILP pairs at all, since the pairs
 177 are indistinguishable in the Weisfeiler–Lehman hierarchy that bounds such
 178 architectures [36]. We therefore do not ask the network for a feasibility ver-
 179 dict; it supplies branching guidance inside a complete search that remains
 180 responsible for soundness.

181 *Graph representations..* We use graph attention [18] over the bipartite variable–
 182 constraint graph, the standard encoding for MILP in this literature [7].

183 *Reinforcement learning for branching..* Imitation targets an expert’s rank-
 184 ing, not the search cost the ranking is meant to reduce. FMSTS [20] uses
 185 a node’s subtree size as an observable Q-value, shows that under depth-
 186 first search minimizing every subtree minimizes the whole tree, and learns
 187 from scratch; tree MDPs [21] formalize that per-subtree credit assignment;
 188 retro branching [31] learns from retrospectively extracted subtree trajecto-
 189 ries; SORREL [25] initializes from suboptimal demonstrations; Neuro# [34]
 190 trains a graph network on the residual formula of a #SAT solver by evolution
 191 strategies to minimize branching decisions. In constraint programming, value
 192 ordering has been learned from solutions [24] and by deep Q-learning [23],
 193 and impact-based search [33] prefers the variable whose assignment shrinks
 194 the search space most. Our fine-tuning stage (Section 3.8) uses the same
 195 per-subtree credit with a plain policy-gradient update [19], and differs in its
 196 starting point: the policy is initialized at a supervised rule derived from exact
 197 conditional marginals, and the marginal network is frozen, so the only thing
 198 the cost signal can change is which variable is branched on. The classical

199 statement of the objective is the fail-first principle [22]: choose the variable
 200 whose wrong branch is refuted soonest.

201 2.4. Amortized optimization

202 The distinction organizing our results predates its use in combinatorial
 203 optimization. Gershman and Goodman [8] introduced *amortized inference*:
 204 when many related queries must be answered, solving each from scratch is
 205 wasteful and the question becomes how to reuse computation across queries.
 206 Amos [3] develops the same idea as *amortized optimization* — using learning
 207 to predict solutions by exploiting shared structure across similar instances
 208 — and surveys its appearance in variational inference, meta-learning and
 209 control. Millidge [13] states the trade-off in the form most relevant here:
 210 *direct* optimization spends computation on the instance at hand and pays
 211 that cost at every invocation, whereas *amortized* optimization converts the
 212 problem into supervised learning with cheap inference but requires a corpus
 213 of solved instances up front and is bounded by the approximator’s gener-
 214 alization; the productive arrangement is hybrid, with amortized predictions
 215 steering a direct optimizer.

216 Our system instantiates that hybrid precisely. The direct method is LP-
 217 probing, sound and costly. The amortized method is the network, cheap and
 218 unsound. The hybrid rule is that the network steers branching while probing
 219 and propagation retain exclusive authority over variable fixing (Section 3.7).

220 3. Proposed Method

221 3.1. Setting and notation

222 For a partial assignment fixing variables $F \subseteq \{1, \dots, n\}$ to values $\mathbf{v} \in$
 223 $\{0, 1\}^{|F|}$, write

$$\mathcal{S}(F, \mathbf{v}) = \{\mathbf{x} \in \mathcal{S} : \mathbf{x}_F = \mathbf{v}\} \quad (3)$$

224 for the solutions consistent with it, and let $K = \{1, \dots, n\} \setminus F$ index the free
 225 variables. The *conditional marginal* of a free variable $j \in K$ is

$$p_j(F, \mathbf{v}) = \frac{|\{\mathbf{x} \in \mathcal{S}(F, \mathbf{v}) : x_j = 1\}|}{|\mathcal{S}(F, \mathbf{v})|}. \quad (4)$$

226 *3.2. System overview*

227 The deployed solver is a cascade of a lattice stage and a guided complete
228 search:

$$\text{propagation} \rightarrow \text{LP infeasibility rule} \rightarrow \text{kernel pump} \rightarrow \text{AHL/BKZ} \rightarrow \text{guided backtracking search} \quad (5)$$

229 The first three stages are cheap and sound. AHL/BKZ attempts a certificate-
230 producing lattice reduction with a size-dependent block parameter; when
231 it succeeds it returns a verified solution outright. The learned component
232 appears only in the final stage, and only in its branching decisions. This
233 placement matters for the evaluation: because the lattice stage is independent
234 of the branching heuristic, every heuristic under comparison receives exactly
235 the same residual set of instances, so conditioning on that residual describes
236 the algorithm rather than selecting favourable cases.

237 *3.3. Conditioning is reduction*

238 **Proposition 1.** *Let A_F and A_K denote the column submatrices of A indexed*
239 *by F and K . Then*

$$\mathcal{S}(F, \mathbf{v}) \cong \{\mathbf{y} \in \{0, 1\}^{|K|} : A_K \mathbf{y} = \mathbf{b} - A_F \mathbf{v}\}, \quad (6)$$

240 *the bijection being $\mathbf{x} \mapsto \mathbf{x}_K$. Consequently the conditional marginals (4)*
241 *of the original instance are the ordinary marginals of the reduced instance*
242 *$(A_K, \mathbf{b} - A_F \mathbf{v})$.*

243 *Proof.* $A\mathbf{x} = A_F \mathbf{x}_F + A_K \mathbf{x}_K$. Fixing $\mathbf{x}_F = \mathbf{v}$ gives $A\mathbf{x} = \mathbf{b} \iff A_K \mathbf{x}_K =$
244 $\mathbf{b} - A_F \mathbf{v}$, and $\mathbf{x} \mapsto \mathbf{x}_K$ is a bijection between the two sets since $\mathbf{x}_F$ is deter-
245 mined. The marginal statement follows by counting. $\square$

246 The identity is elementary but it is what makes in-tree supervision practi-
247 cal. A depth- d search node is not a special object requiring its own encoding;
248 it is an instance of the same problem on $n - d$ variables. One network, one
249 feature extractor and one label generator therefore cover the whole tree, and
250 training data at depth d is produced by substitution rather than by instru-
251 menting a solver.

252 Two implementation details preserve correctness. Rows that become all-
253 zero after substitution carry no constraint and are removed; retaining them
254 causes linear-programming backends to reject the model. And the free bound
255 condition $0 \leq \mathbf{b} - A_F \mathbf{v} \leq A_K \mathbf{1}$ is checked before any relaxation is solved,
256 since it detects a large fraction of dead nodes at negligible cost.

257 3.4. Which predictions can cost anything

258 At a node with surviving solution set $\mathcal{S}(F, \mathbf{v}) \neq \emptyset$, partition the free
259 variables by their conditional marginal:

$$\text{DET} = \{j \in K : p_j \in \{0, 1\}\}, \quad \text{OPEN} = \{j \in K : 0 < p_j < 1\}. \quad (7)$$

260 The distinction is operational rather than statistical:

- 261 • For $j \in \text{OPEN}$, both children of a branch on x_j contain solutions. Either
262 choice keeps a solution reachable, so *no branching decision on a OPEN*
263 *variable can cost anything*.
- 264 • For $j \in \text{DET}$, one child is empty. Branching against the forced value
265 guarantees that the subtree is exhausted without a solution, costing a
266 full backtrack.

267 This gives the conditional marginal its operational reading. It is not a
268 “probability of being correct” about some designated solution, but the *frac-*
269 *tion of the surviving solution set retained* by an assignment; and the only
270 errors that matter are those on variables where that fraction is 0 or 1.

271 Two consequences follow, both of which the experiments confirm. First,
272 an aggregate accuracy over all free variables is a poor objective, because it
273 mixes decisions that can cost a backtrack with decisions that cannot; the
274 relevant quantity is accuracy on DET. Second, since $|\mathcal{S}(F, \mathbf{v})|$ shrinks as the
275 search descends, the DET share grows with depth. This alone does not say
276 *where* a learned predictor is useful, because a DET variable that the relaxation
277 already rounds correctly needs no predictor. The useful region is where DET
278 variables are plentiful *and* the relaxation is weak, and Section 4.9 shows that
279 this is governed by constraint density $m/|K|$ rather than by depth as such:
280 on small instances the two coincide, on larger ones they do not.

281 3.5. Exact conditional marginals as training targets

282 For instances small enough to enumerate, $\mathcal{S}$ is obtained exactly with a
283 constraint solver in all-solutions mode, and (4) is computed by counting.
284 Training minimizes the cross-entropy to the resulting Bernoulli targets,

$$\mathcal{L} = -\frac{1}{|K|} \sum_{j \in K} \left[p_j \log \hat{p}_j + (1 - p_j) \log(1 - \hat{p}_j) \right], \quad (8)$$

285 which is minimized exactly at $\hat{p}_j = p_j$. Note that the same loss with a planted
 286 solution in place of p_j — the conventional choice — is a noisy surrogate
 287 whenever $|\mathcal{S}| > 1$, since the planted solution is one draw from the posterior;
 288 on our instances a Bayes-optimal predictor agrees with it only about two
 289 thirds of the time.

290 *Enumeration correctness.* A subtlety here is easy to get wrong and silently
 291 corrupts every label derived from it. Constraint solvers report distinct sta-
 292 tuses for “search exhausted” and “time limit reached with solutions already
 293 collected.” Only the former certifies that $\mathcal{S}$ is complete; accepting the latter
 294 substitutes a truncated solution set, which biases every marginal computed
 295 from it. In our implementation only **OPTIMAL** and **INFEASIBLE** are accepted
 296 and timed-out instances are discarded rather than averaged in. A useful
 297 diagnostic is that mean enumeration time equal to the time limit indicates
 298 truncated instances being counted as complete.

299 *Sampling reachable states.* A depth- d training state is built by drawing d
 300 uniformly from $\{0, \dots, 12\}$, taking a prefix of an LP-confidence variable or-
 301 dering, and fixing those variables to the values of a randomly chosen *actual*
 302 *solution*. Fixing arbitrary values would produce states with $\mathcal{S}(F, \mathbf{v}) = \emptyset$; such
 303 states are dead nodes that propagation prunes immediately and on which the
 304 conditional marginal is undefined, so they are not states a branching heuris-
 305 tic is ever asked about. The depth range is a design choice we did not tune;
 306 on 10×25 instances it yields training states with 13–25 free variables, and
 307 Section 4.9 examines what happens when the deployed search queries states
 308 outside that range, which at 21×60 is most of them.

309 3.6. Network

310 An instance is encoded as a bipartite graph with n variable nodes, m
 311 constraint nodes, and an edge (j, i) whenever $a_{ij} = 1$. Let $\tilde{\mathbf{x}}$ be the relaxation
 312 solution of the current (reduced) instance and $r_i = \sum_j a_{ij}$ the row sums.
 313 Node features are

$$\text{variable } j : (\tilde{x}_j, \deg(j)/m, |\tilde{x}_j - \text{round}(\tilde{x}_j)|), \quad (9)$$

$$\text{constraint } i : (b_i/r_i, r_i/n, (b_i - A_i\tilde{\mathbf{x}})/r_i), \quad (10)$$

314 i.e. relaxation value, column degree and fractionality for variables; tight-
 315 ness, scaled row size and relaxation residual for constraints. All are size-
 316 normalized, which is what permits a single parameter set to apply at any
 317 (m, n) .

Both node types are embedded to $h = 64$ dimensions by a linear map and refined by $L = 4$ rounds of bidirectional message passing with graph attention [18]:

$$\mathbf{h}^{\text{con}} \leftarrow \text{LN}(\text{ELU}(\text{GAT}(\mathbf{h}^{\text{var}} \rightarrow \mathbf{h}^{\text{con}})) + \mathbf{h}^{\text{con}}), \quad (11)$$

$$\mathbf{h}^{\text{var}} \leftarrow \text{LN}(\text{ELU}(\text{GAT}(\mathbf{h}^{\text{con}} \rightarrow \mathbf{h}^{\text{var}})) + \mathbf{h}^{\text{var}}), \quad (12)$$

followed by a two-layer head emitting one logit per variable, $\hat{p}_j = \sigma(z_j)$. The model has 21,761 parameters.

The architecture is deliberately small and single-headed. Its only output is the quantity the objective (8) defines; there are no auxiliary value or feasibility heads, because the search draws soundness from its symbolic components and asks the network for nothing else.

3.7. Inference: guidance, not commitment

The network is consulted once per node. Branching selects

$$j^* = \arg \max_{j \in K} |\hat{p}_j - \tfrac{1}{2}|, \quad \text{first value} = \mathbb{K}[\hat{p}_{j^*} \geq \tfrac{1}{2}], \quad (13)$$

that is, the variable the network is most confident about, tried at its predicted value. The search remains complete and a wrong prediction costs backtracking only.

Variables are *fixed* only by sound deduction. Two mechanisms are available: constraint propagation, and LP-probing, which fixes x_j when assigning the opposite value makes even the relaxation infeasible. The division of labour follows directly from the measurements in Section 4.5: where a wrong answer costs time, the cheap unsound predictor is preferable; where a wrong answer costs correctness, the expensive sound procedure is required and cannot be replaced.

Algorithm 1 states the resulting search.

3.8. Cost-aware fine-tuning of the branching policy

The rule of Algorithm 1 branches on the variable the network is surest about. That maximizes the probability that the first child contains a solution; it says nothing about how much the search pays when the first child does not, nor about how much smaller the children are. Both quantities are properties of the *subtree*, and a search produces them as a by-product: the cost of a branching decision at node u is the number of nodes expanded below u until

Algorithm 1 Guided backtracking search for a BLS

Input: instance $(A, \mathbf{b})$, guidance model f_θ , time budget T
stack $\leftarrow [(A, \mathbf{b})]$
while stack $\neq \emptyset$ **and** elapsed $< T$:
 $(A', \mathbf{b}') \leftarrow \text{pop}(\text{stack})$
 if $\mathbf{b}' < \mathbf{0}$ **or** $\mathbf{b}' > A'\mathbf{1}$: **continue** // *free bound check*
 drop all-zero rows of A' ; **if** none remain: **return** any completion
 propagate; **if** contradiction: **continue**
 if propagation determined a set D of variables:
 push reduce($A', \mathbf{b}', D$); **continue** // *Proposition 1*
 if all variables assigned: **return** solution if verified, else **continue**
 $\hat{\mathbf{p}} \leftarrow f_\theta(A', \mathbf{b}')$ // *one forward pass, all variables*
 $j^* \leftarrow \arg \max_j |\hat{p}_j - 1/2|$; $v \leftarrow \mathbb{K}[\hat{p}_{j^*} \geq 1/2]$
 push reduce($A', \mathbf{b}', \{j^* = 1-v\}$), then reduce($A', \mathbf{b}', \{j^* = v\}$)
return unresolved

347 a solution is found there or the subtree is exhausted, and depth-first search
348 returns it exactly when it backtracks out of u . We use this cost to fine-tune
349 which variable is branched on.

350 *Decision process..* A state is the reduced instance $(A', \mathbf{b}')$ after propagation;
351 an action is a free variable j ; the first value tried remains $\mathbb{K}[\hat{p}_j \geq 1/2]$ from
352 the marginal network, which is not modified. The return of the action at u
353 is $-\log c(u)$ with $c(u)$ the subtree cost. Because every decision has its own
354 subtree, credit is assigned per decision rather than per episode, which is what
355 made the sparse-reward formulations we tried earlier unlearnable.

356 *Policy and value..* On the frozen MarginalNet trunk we add two heads. The
357 policy logit is

$$g_j(s) = \log |\hat{p}_j - \tfrac{1}{2}| + h_\phi(\mathbf{h}_j), \quad (14)$$

358 where h_ϕ is a two-layer network on the variable embedding $\mathbf{h}_j$ whose last
359 layer is initialized to zero, so that $\pi(j | s) \propto \exp(g_j/\tau)$ starts as a softening
360 of the supervised rule and its argmax is exactly Algorithm 1. The value head
361 $V_\psi(s)$ reads a pooled graph embedding together with $\log |K|$ and $\log m$ and
362 predicts $\log c$.

363 *Update..* Each epoch samples training states, rolls the sampled policy ($\tau = 1$)
364 out to a solution, and records every decision with its exact cost. With ad-

365 vantage $\alpha(u) = V_\psi(s_u) - \log c(u)$, standardized within the epoch, the update
 366 is

$$\mathcal{L} = \sum_u \left[-\alpha(u) \log \pi(j_u | s_u) - \lambda \mathcal{H}(\pi(\cdot | s_u)) + (V_\psi(s_u) - \log c(u))^2 \right], \quad (15)$$

367 with entropy weight $\lambda = 0.01$. Only ϕ and ψ are trained; the trunk and the
 368 marginal head are fixed, so the marginal used for the first value and for the
 369 analysis of Section 3.4 is unchanged. The critic is pre-trained on decisions
 370 recorded under the supervised policy before the actor is updated, and its
 371 held-out rank correlation with the true cost is used as a gate: if the baseline
 372 cannot rank costs, the advantage is noise and fine-tuning is not attempted.
 373 Epoch selection uses the total subtree cost on held-out training-distribution
 374 states, never the evaluation instances. At deployment the policy is used
 375 greedily: $j^\star = \arg \max_j g_j(s)$.

376 3.9. Cost model

377 The comparison that motivates the design is per-node cost for a whole-
 378 instance verdict, and its structure is asymptotic rather than a constant factor.
 379 Let n_c be the number of free variables at a node. The network solves *one*
 380 relaxation to build its features and performs *one* forward pass, so its LP
 381 count is $O(1)$ in n_c . LP-probing solves one relaxation per candidate variable,
 382 $O(n_c)$. The ratio between them therefore grows linearly in instance size, and
 383 any single number quoted for it is meaningful only together with the n_c at
 384 which it was measured.

385 This also means the comparison is sensitive to how the LP backend is
 386 implemented, in a way that cuts against the learned method. A backend
 387 that rebuilds the model for every call inflates probing’s cost by its full $O(n_c)$
 388 factor while charging the network only once, exaggerating the advantage. A
 389 correctly engineered branch-and-bound keeps one relaxation and re-optimizes
 390 from the parent basis after a bound change, which is an order of magnitude
 391 cheaper per probe. Section 4.5 measures both methods against a warm-
 392 started backend for this reason, and reports the ratio as a function of n_c
 393 rather than as a single figure.

394 4. Experiments

395 4.1. Instances and control of solution multiplicity

396 Instances are generated by the procedure $\text{GENHARD}(m, n, K)$: draw a
 397 random 0/1 matrix A of density $1/2$; plant $K = 20$ candidate solutions
 398 $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(K)}$ with $\mathbf{b}^{(k)} = A\mathbf{x}^{(k)}$; retain the one maximizing *vertex spread*, the
 399 mean pairwise L_2 distance among relaxation vertices obtained from random
 400 objective directions. Maximizing vertex spread enlarges $\tilde{\mathcal{S}}$ directly and is
 401 the mechanism that makes market-split instances hard for relaxation-based
 402 methods (Section 2). Uniqueness of the planted solution is not enforced.

403 *Why $|\mathcal{S}|$ must be controlled across sizes..* Solution multiplicity changes the
 404 *character* of the prediction task, not only its difficulty. When $|\mathcal{S}| = 1$ every
 405 free variable is DET and predicting well is tantamount to solving; when $|\mathcal{S}|$ is
 406 large most variables are OPEN and the task is to estimate a diffuse posterior.
 407 Comparing sizes without controlling $|\mathcal{S}|$ therefore confounds size with task
 408 identity. Since $|\mathcal{S}|$ falls as m rises at fixed n , we select m per size to match
 409 it.

Size	m	m/n	median $ \mathcal{S} $	root ceiling	conditional ceiling
10×25	10	0.40	23	70.4%	87.0%
18×50	18	0.36	19	71.3%	87.2%
21×60	21	0.35	10	—	—

Table 1: Sizes used, with solution multiplicity matched. Ceilings (the Bayes-optimal per-variable accuracies of Section 4.3) agree to within one point between the two training sizes, so size is the variable that changes.

410 *Limits of the control..* Matching $|\mathcal{S}|$ forces m/n to vary ($0.40 \rightarrow 0.35$); at
 411 fixed n the two cannot be held simultaneously and we prioritize $|\mathcal{S}|$ because
 412 it determines task identity. The control is also granular: at $n = 60$ the
 413 median $|\mathcal{S}|$ steps $42 \rightarrow 12 \rightarrow 2$ across $m = 20, 21, 22$, so $m = 21$ is the best
 414 available integer and the realized evaluation set has median $|\mathcal{S}| = 10$ with
 415 $27/30$ verified.

416 *Where enumeration stops..* At $n = 70$ no enumeration completed within 120 s
 417 at the m giving the target multiplicity (0/10). At $n = 100$ the obstruction
 418 is more basic: a constraint solver finds *zero* solutions within 20 s at $m \in$

419 $\{28, 32, 36, 40\}$ despite a planted one. At that size the bottleneck is finding
 420 any solution rather than choosing branches well, so a branching comparison
 421 there would not be informative. We therefore evaluate at $n \leq 60$ and report
 422 this boundary as a limitation.

423 *4.2. Implementation and protocol*

424 *Software..* Relaxations used for training features and for the original search
 425 loop are solved with SCIP via PySCIPOpt; the warm-started relaxations
 426 of Sections 4.5 and 4.6 use Gurobi 13 (dual simplex, bounds changed in
 427 place), which requires a licence — HiGHS exposes an equivalent incremental
 428 interface and should serve as a free substitute, though we have not verified
 429 the timings with it. Enumeration and the CP-SAT baseline use OR-Tools;
 430 lattice reduction uses fpylll (LLL and BKZ). The network is implemented in
 431 PyTorch with PyTorch Geometric (GATConv). One property of the relaxation
 432 deserves note: it carries no objective, so its solution is an arbitrary vertex
 433 of the relaxation polytope and different LP codes return different vertices.
 434 Both branching arms consume the same vertex within an experiment, so
 435 comparisons are fair, but absolute node counts differ between the SCIP- and
 436 Gurobi-backed runs.

437 *Training..* Adam, learning rate 10^{-3} , 20 epochs over 10,000 in-tree states,
 438 batch size 1 (instances vary in size), gradient-norm clipping at 1.0, loss (8)
 439 with exact marginal targets. Training states are drawn at depths 0–12 uni-
 440 formly, 6 states per source instance, from a pool of 1,760 instances; evaluation
 441 uses 440 instances from a disjoint pool. Training takes approximately 20 min-
 442 utes on one CPU core. *No hyperparameter was tuned:* width, depth, learning
 443 rate, epoch count and the training depth range are the values first tried, and
 444 there is no validation split. Reported figures are therefore not contaminated
 445 by selection on the test set, but neither is any of these settings claimed to be
 446 optimal.

447 *Threading..* All timing-sensitive components run single-threaded (`torch.set_`
 448 `num_threads(1)`). On graphs of this size multi-threaded BLAS is counter-
 449 productive: we measure 11.3ms per forward pass at 8 threads against 0.9ms
 450 at one.

451 *Search evaluation..* Each arm is run on identical instances with a hard per-
 452 instance wall-clock budget enforced by subprocess termination, since lattice

reduction executes inside a blocking native call that in-process timers cannot interrupt. Instances run at concurrency 6 on a 20-core machine. This low concurrency is deliberate: wall-clock-budgeted comparisons are sensitive to CPU contention, and at higher concurrency we have observed instances solvable in 14–15s standalone exceeding a 20s cap.

Statistics. Solve-rate differences between arms on the same instances are tested with McNemar’s exact test; per-instance speed differences with a two-sided sign test. We report both rather than a single aggregate because they can have different power. The main comparison at 21×60 is replicated over three independent seeds and we report the spread; smaller-size and transfer results are single-run and marked as such. Seeds are set explicitly: an earlier version derived them from `hash()` of the instance filename, which Python randomizes per process, so those runs had no recorded provenance.

Evaluation sets. Two 21×60 sets are used: the original 30 instances (generator seeds 700000–700029) and a 100-instance set (seeds 700000–700099) that contains them, built after the 30-instance results had been reported; $|\mathcal{S}|$ is verified for 92/100, median 11. All numbers on the 100-instance set are lattice-off, three seeds. The training pool for the experiments of Sections 4.10 and 4.11 is a separate set of 500 instances at 21×60 (seeds 800000–800499, split 400/100) whose overlap with every evaluation set is zero by hash of $(A, \mathbf{b})$; the 100 held-out instances of that pool supply the states on which fine-tuning epochs are selected.

4.3. Prediction quality against the Bayes ceiling

Because $\mathcal{S}$ is enumerated exactly, the Bayes-optimal per-variable accuracy is computable:

$$\text{ceil} = \frac{1}{|K|} \sum_{j \in K} \max(p_j, 1 - p_j), \quad (16)$$

attained by answering the majority value of the surviving solutions. No deterministic predictor of $\mathbf{x}^*$ from $(A, \mathbf{b})$ can exceed it. Table 2 evaluates predictors against it on 400 held-out 10×25 instances with $\mathcal{S}$ fully enumerated.

The network comes within 4.1 points of the ceiling overall and 6.6 on the top-3 subset, and it recovers most of the gap between relaxation rounding and the ceiling ($60.9 \rightarrow 66.7$ against a ceiling of 70.8). Two preliminary experiments on an earlier pool of 10×25 instances, reported for completeness

Predictor	all variables	top-3	L_1 to true marginal
Bayes ceiling	70.8%	93.2%	0.0000
Proposed network	66.7%	86.6%	0.1216
LP relaxation rounding	60.9%	65.8%	—

Table 2: Root-level prediction against the exact posterior on the 182 root states of the held-out 10×25 evaluation pool. The network is the same checkpoint used in every search experiment below. “top-3” restricts to the three variables the predictor is most confident about.

486 but not from the checkpoint evaluated here, bear on design choices: a 2.11M-
487 parameter network with four output heads (selection, assignment, value and
488 feasibility) scored 60.7% on root states, statistically indistinguishable from
489 relaxation rounding at 60.6%, so capacity is not the binding constraint and
490 a single-purpose head is preferable; and training the single-head architecture
491 on exact marginals rather than the planted solution improved L_1 to the
492 posterior from 0.1261 to 0.0977 while leaving argmax accuracy unchanged,
493 which matters because the analysis of Section 3.4 consumes the marginal
494 rather than its argmax.

495 *What the aggregate conceals..* Table 3 decomposes the same root variables
496 by the partition of Section 3.4.

Root variables	count	share	ceiling on subset
DET (branching against it empties the subtree)	38	6.3%	100.0%
OPEN (either branch reaches a solution)	562	93.7%	67.2%
all variables	600	100%	69.3%

Table 3: Root-level decomposition, 24 instances of 10×25 with $\mathcal{S}$ enumerated. The aggregate is 94% composed of variables on which no branching decision can cost a backtrack.

497 At the root, 93.7% of variables are OPEN. The aggregate ceiling of 69.9%
498 is therefore dominated by predictions that cannot affect the search, and on
499 the 6.3% that can, the information is already complete (100%). Root-level
500 accuracy is consequently a weak objective for a branching heuristic, and the
501 operative question is not whether the information exists but what detects it
502 — taken up next.

503 4.4. *Depth, and where the residual lies*

504 Using Proposition 1 we measure, at each depth along a prefix drawn
 505 from an actual solution, the conditional ceiling and the fraction of remaining
 506 variables that propagation forces (Table 4).

depth	surviving $ \mathcal{S} $	conditional ceiling	proposed network	LP rounding
0	25.3	70.8%	66.7%	60.9%
4	4.2	82.3%	72.8%	68.5%
7	1.8	91.5%	81.2%	77.8%
8	1.4	94.4%	84.3%	80.9%
12	1.1	98.8%	96.8%	95.5%
all depths	—	87.2%	80.0%	76.5%

Table 4: Conditional information against depth on the held-out in-tree states of the 10×25 evaluation pool ($n = 25$), LP-confidence variable ordering. Same checkpoint as Table 2.

507 The consequential movement is in composition rather than level: the
 508 DET share rises from 6.3% at the root to 87.9% at depth 8 (Table 6). At the
 509 root almost every branching decision is inconsequential because both children
 510 contain solutions; by depth 8 almost every decision empties a subtree when
 511 wrong. The network’s margin over relaxation rounding is 3–6 points through-
 512 out and does not vanish at depth, which is the first indication that what it
 513 captures is not simply depth-dependent; Section 4.9 makes this precise.

514 *Does in-tree supervision matter, controlling for everything else?* A compari-
 515 son between a root-trained and an in-tree-trained network is only informative
 516 if nothing else differs. Table 5 holds the instance pool, the label type (exact
 517 marginals), the architecture, the optimizer and the number of training states
 518 (1,760 each) fixed, and varies only whether the states are root states or in-
 519 tree states drawn at depths 0–12. Both are evaluated on the same held-out
 520 in-tree states.

521 In-tree supervision is worth 4.4 points overall. The more telling row is the
 522 root-trained network itself: at 75.4% it is *below* relaxation rounding on in-
 523 tree states, so a network trained on root states does not merely underperform
 524 in the tree, it is worse than the classical heuristic it is meant to replace. The
 525 shift from root to interior states is not a distribution the root-trained network
 526 generalizes across, and in-tree training is what makes the learned guidance
 527 better than the baseline at all.

depth	conditional ceiling	root-trained	in-tree-trained	LP rounding
0	70.7%	67.8%	66.5%	60.9%
4	82.1%	71.4%	72.4%	68.3%
6	89.7%	74.6%	78.2%	74.0%
8	94.4%	77.4%	84.2%	81.3%
10	97.1%	81.7%	91.4%	89.4%
12	98.8%	85.8%	96.5%	95.3%
all depths	87.2%	75.4%	79.8%	76.4%

Table 5: Controlled comparison of training-state distribution. Same pool, labels, architecture, optimizer and state count; only the depth at which states are sampled differs. Evaluated on held-out in-tree states of the 10×25 pool.

528 Table 6 decomposes the residual gap at these depths and asks what each
529 method detects.

depth	% DET	network on DET	propagation	LP integrality	LP-probing	network on OPEN
5	61.9%	89.7%	1.2%	67.9%	51.6%	54.4%
6	74.2%	88.2%	1.6%	70.4%	68.0%	52.5%
7	79.9%	88.8%	3.7%	73.3%	72.8%	51.2%
8	87.9%	89.4%	5.3%	76.9%	79.8%	53.8%

Table 6: Residual decomposition on the same states as Table 4. Columns 4–6 are recall on the ground-truth DET set (LP-probing measured with the warm-started backend of Section 4.5); columns 3 and 7 are the network’s argmax accuracy on each subset.

530 The gap lies entirely on DET variables: at depth 7, $0.80 \times 11.2 \approx 9$ points,
531 matching the measured 10.3-point shortfall against the conditional ceiling.
532 Performance near 50% on OPEN variables is correct behaviour rather than
533 a defect, since those variables are genuinely undetermined. The striking
534 figure is that unit propagation detects only 1.2–5.3% of logically determined
535 variables: the information is present and derivable but not reachable by the
536 cheapest sound rule. LP-probing recovers 52–80% and LP integrality 68–
537 77%.

538 4.5. Cost per node

539 Both methods are charged everything they must pay to reach a whole-
540 instance verdict: the network one relaxation for its features plus one forward

pass, probing one relaxation per free variable. Both use the same warm-started backend, a single model whose variable bounds are changed in place and re-optimized from the parent basis (Gurobi dual simplex), which is how a branch-and-bound implementation reuses its relaxation. Timings are single-threaded.

Instances	n	network	LP-probing	ratio	probing recall
10×25	25	3.22 ms	5.60 ms	$1.7\times$	—
18×50	50	2.33 ms	11.84 ms	$5.1\times$	—
21×60	60	2.57 ms	18.39 ms	$7.2\times$	—
40×100	100	4.02 ms	62.46 ms	$15.5\times$	—
60×150	150	4.45 ms	79.06 ms	$17.8\times$	—
<i>in-tree states of 10×25 instances, by depth:</i>					
depth 5	20	2.44 ms	4.02 ms	$1.6\times$	51.6%
depth 6	19	2.39 ms	3.98 ms	$1.7\times$	68.0%
depth 7	18	2.35 ms	3.43 ms	$1.5\times$	72.8%
depth 8	17	2.34 ms	3.13 ms	$1.3\times$	79.8%

Table 7: Per-node cost for a whole-instance verdict, warm-started LP backend for both methods, single-threaded. The network’s accuracy on DET variables at depths 5–8 is 85.9–92.7% against probing’s recall in the last column.

The ratio behaves as the cost model predicts: it is negligible at $n = 25$ and reaches $17.8\times$ at $n = 150$, tracking the $O(n_c)$ versus $O(1)$ difference in relaxation count. Two consequences follow, and they point in opposite directions.

The advantage is real but asymptotic, and small where we can search.. At the sizes where our search actually operates ($n = 25$ – 60) the ratio is 1.3 – $7.2\times$. At the depths where prediction matters most (5–8 of a 10×25 instance, $n_c = 17$ – 20) it is 1.3 – $1.7\times$. A practitioner choosing between the two at that scale should prefer probing, which is *sound*: where its test fires the conclusion is a proof, whereas the network never offers a guarantee. The learned alternative becomes compelling only at $n \gtrsim 100$, and Section 4.1 shows that search itself fails there for every strategy we tried.

The measurement is sensitive to baseline engineering.. An earlier version of this comparison used a backend that rebuilt the LP model on every call

and, additionally, excluded the network’s own feature relaxation from its timer. That combination reported 19–23 $\times$ at depths 5–8, where the correctly measured value is 1.3–1.7 $\times$ — an overstatement of more than an order of magnitude, arising entirely from how the two sides were instrumented rather than from what they compute. We report the corrected figures and note the general hazard: in any cost comparison against a classical procedure, the classical side must be implemented as a practitioner would implement it, and the learned side must be charged for its own preprocessing.

4.6. Search performance

We compare branching strategies under matched wall-clock budgets with the lattice stage disabled, isolating branching quality. LP selects the variable with the most integral relaxation value and tries that value first; MODEL uses the proposed network. Both run in the same loop: one warm-started relaxation per node (Section 4.2), identical propagation, identical value ordering, so the arms differ only in which variable they branch on and which value they try first. Each configuration is run three times with different seeds; the depth-first search itself is deterministic given the relaxation vertex, so seeds change timing but not node counts.

Size (budget)	strategy	solve rate	median nodes	median time	nodes/s
10×25 (60 s)	LP	100 $\pm$ 0.0%	46	0.03 $\pm$ 0.00 s	1,767
	MODEL	100 $\pm$ 0.0%	24	0.07 $\pm$ 0.01 s	376
18×50 (300 s)	LP	100 $\pm$ 0.0%	4,224	2.43 $\pm$ 0.33 s	2,066
	MODEL	100 $\pm$ 0.0%	1,085	1.82 $\pm$ 0.37 s	731
21×60 (600 s)	LP	100 $\pm$ 0.0%	82,366	51.5 $\pm$ 2.0 s	1,627
	MODEL	100 $\pm$ 0.0%	8,657	14.6 $\pm$ 0.7 s	619

Table 8: Branching strategy comparison on the warm-started loop, lattice stage disabled, 30 instances per cell, mean $\pm$ standard deviation over three seeds of the per-seed median. Node counts are seed-invariant.

Three points. First, the learned guidance produces a smaller tree at every size — 2.0 $\times$, 3.9 $\times$ and 9.5 $\times$ fewer nodes than LP — and the ratio grows with size, as the analysis of Section 4.9 predicts: the states where the network separates from relaxation rounding are the large, weakly constrained ones near the top of a big tree, and 10×25 has few of them. Second, a smaller tree buys time only once the tree is large enough to pay for the forward pass.

At 10×25 the MODEL arm is *slower* (0.34–0.43 $\times$, faster on at most 1/30 instances), because a 24-node tree costs less than the network evaluations that shrink it; its per-node cost is 2.6–4.7 $\times$ that of LP at every size, which is the forward pass of Section 4.5. At 18×50 the median per-instance ratio is 1.30–1.37 $\times$ in the network’s favour but not significant (faster on 18/30 in every seed, sign test $p = 0.36$). Third, at 21×60 the advantage is clear and replicates: the median per-instance ratio is 2.57 $\times$, 3.03 $\times$ and 2.96 $\times$ across seeds, the network is faster on exactly 23/30 instances in each ($p = 0.005$), and both arms solve every instance within budget.

Two further strategies were measured on an earlier version of this loop that rebuilt the relaxation at every node. RANDOM branching solved 0/30 at 20×50 within 300s, and Section 4.12 shows it closes none of the 21×60 residual instances on the present loop either. LP-PROBE, which probes every free variable and soundly fixes the forced ones before branching as LP, produced by far the smallest trees (51 $\times$ fewer nodes than LP at 20×50) and was the slowest arm, processing 6 nodes/s against 376: node count and wall-clock rank the methods in opposite orders, and a tree-size convention would select the worst of them. We have not repeated LP-PROBE on the warm-started loop; Section 4.5 measures its per-node cost there directly.

Replication on 100 instances.. Thirty instances leave the sign test with little power and the median with wide uncertainty, so the 21×60 comparison was repeated on the 100-instance set (Table 9). The direction and significance hold — MODEL is faster on 70–71/100 instances ($p < 0.001$) with 5.5 $\times$ fewer nodes — but the margins are smaller than on the original thirty: 2.0–2.3 $\times$ in time against 2.6–3.0 $\times$, and 5.5 $\times$ in nodes against 9.5 $\times$. The original thirty are the first thirty seeds of the larger set, and on them the supervised policy happens to do unusually well; we treat the 100-instance figures as the estimate and the 30-instance ones as the favourable end of its sampling variation.

strategy	solve rate	median nodes	median time	nodes/s	vs. LP (per instance)
LP	100 $\pm$ 0.0%	79,609	41.2 $\pm$ 1.9s	1,966	—
MODEL	100 $\pm$ 0.0%	13,386	16.5 $\pm$ 0.6s	796	2.0–2.3 $\times$, 70–71/100, $p < 0.001$
MODEL+FT (Section 4.11)	100 $\pm$ 0.0%	5,495	7.6 $\pm$ 0.3s	738	4.8–5.9 $\times$, 81–84/100, $p < 0.001$

Table 9: 21×60 , 100 instances, lattice off, 600s budget, three seeds (mean $\pm$ s.d. of per-seed medians; node counts are seed-invariant). The last row is the cost-aware fine-tuned policy of Section 4.11.

613 We also record what the rebuilt-relaxation loop reported for the 21×60
614 comparison, because an earlier draft of this paper rested on it: 30/30 against
615 25/30 solved, and a $3.29\times$ median speedup. Warm-starting the relaxation
616 removed the solve-rate difference entirely — the five instances LP failed within
617 600 s it now closes in 65–270 s — and reduced the speedup. The solve-rate
618 advantage was an artifact of a slow backend that penalized the arm expanding
619 more nodes, and we do not claim it.

620 4.7. Comparison against complete solvers

621 Table 8 compares branching strategies inside our own search loop, which
622 measures branching quality but says nothing about how the system stands
623 against an off-the-shelf solver. That comparison is reported here, on the same
624 instances, single-threaded, under the same budget.

Size	CP-SAT	SCIP	proposed (MODEL)	CP-SAT advantage
10×25	0.002 s	0.009 s	0.07 s	$35\times$
18×50	0.046 s	0.385 s	1.82 s	$40\times$
21×60	1.159 s	3.687 s	14.6 s	$13\times$
21×60 , 100 inst.	0.792 s	2.590 s	16.5 s / 7.6 s (fine-tuned)	$21\times$ / $10\times$

Table 10: Median solve time; all methods solve every instance at every size. The proposed column is the MODEL arm of Table 8 (mean of per-seed medians); the last row adds the 100-instance set and the fine-tuned policy of Section 4.11. The proposed system does not beat a complete solver at any size tested, and is outperformed by SCIP, whose LP code it uses for training features.

625 The gap is large and we state it plainly: *the contribution of this paper is*
626 *not a solver that competes with CP-SAT*. It is worth asking where the gap
627 comes from, because the answer bears on what the branching heuristic is
628 actually doing. Table 11 decomposes it at 21×60 into tree size and per-node
629 throughput, using CP-SAT’s own branch counter.

630 The learned guidance produces a *smaller* search tree than CP-SAT —
631 $2.2\times$ fewer nodes despite CP-SAT’s clause learning — and loses entirely on
632 per-node cost. Where that cost goes can be measured. Before the loop was
633 warm-started it processed 254 nodes/s, spending 4.34 ms per node rebuilding
634 and solving the relaxation; warm-starting the same LP brings that to 0.08 ms
635 and raises throughput to 619 nodes/s, which is the loop reported throughout
636 Section 4.6. Of the 1.6 ms that remain per node, the network’s forward pass
637 is 1.88 ms at the root and less at deeper nodes where the reduced graph is

	CP-SAT	proposed	ratio
nodes (median)	18,827 branches	8,657	ours 2.2× fewer
throughput	19,782 nodes/s	619 nodes/s	CP-SAT 32× faster
total time (median)	1.065 s	14.6 s	14×

Table 11: Decomposition of the 21×60 gap. CP-SAT additionally learns a median of 5,514 conflict clauses per instance.

638 smaller, and propagation and bookkeeping are under 0.1 ms: the forward
639 pass now accounts for over 90% of per-node cost, and we have not reduced
640 it. We therefore do not claim that the gap is an implementation artifact.
641 What the node counts establish is narrower: the deficit is not in branching
642 quality. Closing it would require at minimum a compiled inference path and,
643 as the next paragraph notes, algorithmic components our search lacks.

644 Two honest caveats attach to this reading. Tree sizes are not perfectly
645 commensurable — CP-SAT’s “branches” include restarts and conflict-driven
646 jumps, and each of its nodes performs more propagation work than ours
647 — so the 2.2× should be read as an order-of-magnitude statement. And
648 a competitive implementation would need more than a faster forward pass:
649 clause learning and restarts are absent from our search, and the last two are
650 what let CP-SAT scale past the sizes where we stop.

651 4.8. Size transfer

652 Because the features are size-normalized and the parameters do not de-
653 pend on (m, n) , one network applies at any size. We test whether that
654 transfers in practice, with $|\mathcal{S}|$ matched as in Table 1.

Training size	Evaluation size	solve rate	median time	median nodes
10×25	10×25	100%	0.07 ± 0.01 s	24
10×25	18×50	100%	1.82 ± 0.37 s	1,085
18×50	18×50	100%	1.48 ± 0.15 s	942
10×25	21×60	100%	14.6 ± 0.7 s	8,657
18×50	21×60	100%	23.7 ± 0.7 s	14,742

Table 12: Transfer on the warm-started loop, lattice stage disabled, 30 instances per row, three seeds. Rows 2–3 and 4–5 each hold the evaluation size fixed and vary only the training size.

At the 21×60 evaluation size the network trained only at 10×25 — a $2.4 \times$ smaller problem — is not worse than the one trained at 18×50 ; if anything it is better, expanding fewer nodes on $18/30$ instances and finishing faster on $18\text{--}19/30$ in each seed, with a median per-instance time ratio of $1.24\text{--}1.31 \times$ in its favour. Neither difference is significant (sign test $p \geq 0.20$), and both solve $30/30$. At 18×50 the picture is the same: the 10×25 network is faster on $16\text{--}17/30$ instances ($p \geq 0.59$) although its median time is slightly higher, the two summaries disagreeing because the paired differences are small against instance-to-instance variation. A previous version of this comparison, on the rebuilt-relaxation loop, gave $15/30$ and a $0.92 \times$ ratio; the direction is not stable across loops and we read the result as no detectable transfer loss rather than as a transfer gain. Root-level prediction agrees: the 10×25 network scores 67.8% , 69.5% and 69.4% at 10×25 , 20×50 and 40×100 , holding its margin over relaxation rounding ($+11.0$, $+7.2$, $+4.7$ points) at every size.

This matters practically because label generation is the binding constraint. Exact marginals require enumeration, which fails beyond $n \approx 60$; transfer means a network trained where labels are obtainable can be deployed where they are not.

4.9. Where the transferred network helps

Section 4.4 located the useful prediction regime at depths 6–8 of a 10×25 instance. Whether that localization survives transfer is a separate question, and it does not: at 21×60 the search queries the network on states with 5–60 free variables, only 19.7% of which fall inside the training range of $13\text{--}25$, and 80.2% of which are larger than any training state. Table 13 evaluates the network and relaxation rounding against exact conditional marginals on 21×60 in-tree states, grouped by the number of free variables, on the subset of states that could be enumerated within 15 s.

The pattern is the reverse of what a depth-based account predicts. Inside the training range the network is indistinguishable from relaxation rounding — both are perfect and their outputs correlate at 1.000 — because at $m/|K| \geq 0.8$ the reduced instance is over-constrained and the relaxation is already integral on every DET variable. The network’s entire margin arises on states it never saw during training, at $50\text{--}60$ free variables, where the relaxation is weakest ($m/|K| \approx 0.35\text{--}0.42$) and its outputs have decorrelated from the relaxation’s. In terms of Section 3.4, this is the region where DET variables are present but the relaxation does not find them.

free vars	in training range	$m/ K $	network on DET	LP on DET	gain	corr(network, LP)
20	yes	1.05	100.0%	100.0%	+0.0	1.000
25	yes	0.84	100.0%	100.0%	+0.0	0.998
30	no	0.70	95.7%	95.0%	+0.7	0.954
40	no	0.53	75.5%	75.2%	+0.4	0.747
50	no	0.42	70.2%	65.0%	+5.3	0.664
60	no	0.35	93.3%	88.9%	+4.4	0.658

Table 13: 21×60 in-tree states by number of free variables, 9–14 enumerable states per row. Gain is the network’s margin over relaxation rounding on DET variables; the last column is the Pearson correlation between the two predictors’ outputs.

Two conclusions follow. The first corrects Section 4.4: depth was a proxy for constraint density that happened to hold on 10×25 , where descending from the root moves $m/|K|$ only from 0.40 to 0.59, and does not hold on 21×60 , where the same descent moves it from 0.35 past 1.0. The useful regime is defined by $m/|K|$, and at deployment size it sits near the *top* of the tree. The second is that the transfer result of Section 4.8 is stronger than a matched-multiplicity comparison alone shows: the network extrapolates to states outside its training range in both size and constraint density, and that extrapolation is where its contribution lives. It is also a limitation, since the training range was not chosen with this in mind (Section 4.14).

4.10. Training on deployment-sized states: a negative result

Section 4.9 located the network’s margin on states with 50–60 free variables, outside the 13–25 range of its training states, and Section 4.14 of an earlier draft listed aligning the two as an obvious untested improvement. We tested it, holding everything but the training-state distribution fixed.

States. From a separate 21×60 pool (Section 4.2) we built in-tree states with 30, 35, ..., 60 free variables by fixing a prefix of the LP-confidence ordering to the planted solution, and labelled them with exact conditional marginals under the enumeration rule of Section 3.5. Every state with at most 55 free variables enumerated within 90 s; at 60 (the root) 291/400 did, and the failures are the instances with large $|\mathcal{S}|$, so root labels are biased toward small solution sets. A property of these states matters for what follows: with the prefix fixed to a solution, the surviving solution set is almost always a singleton once 50 or fewer variables remain (median $|\mathcal{S}(F, \mathbf{v})| = 1$, DET share 97%), so their labels are nearly hard. Three networks were trained with the

717 architecture, optimizer, epochs and loss of Section 4.2: M0, the network used
718 everywhere above (10×25 in-tree states); M1, the 2,691 deployment-range
719 states only; M2, half of each. A pre-diagnostic on 10×25 alone, training
720 on depth ranges 0–6, 6–12 and 0–12 with equal state counts, gave in-tree
721 accuracies of 79.3%, 79.9% and 79.9%: within one size, the depth range does
722 not matter.

723 *Prediction..* Re-running the analysis of Table 13 with the three networks
724 scored on identical states, the gain over relaxation rounding on DET variables
725 at 50 and 60 free variables is +5.4 and +9.7 points for M0, +6.1 and +10.7
726 for M1, +5.9 and +10.6 for M2. Training in the deployment regime buys
727 about one point there.

728 *Search..* It costs far more elsewhere. On the 100-instance set M1 expands
729 a median of 21,337 nodes against 13,386 for M0 ($1.6\times$ more; fewer on only
730 38/100, $p = 0.021$) and is $1.7\text{--}1.8\times$ slower; on the original thirty it expands
731 20,746 against 8,657. M2 sits between (18,558, $p = 0.089$). A switch that
732 routes nodes with ≥ 50 free variables to M1 and the rest to M0 recovers
733 nothing (17,252 nodes, 43/99, $p = 0.23$). The mechanism is visible in the
734 prediction diagnostics: on small, dense interior states (10×25 at depth 6–
735 8, the kind the search spends two thirds of its nodes on) the probability
736 that M1’s most confident variable is DET and predicted correctly falls from
737 94.4% to 77.6%, and the rank correlation between its confidence and the true
738 margin from 0.36 to 0.12, while on the deployment-range states it improves
739 ($77.2 \rightarrow 86.0\%$). Nodes are where the interior is, and a network that is
740 slightly better near the top of the tree and much worse inside it makes the
741 search slower. We conclude that the training-state size distribution is not the
742 lever, and that the residual identified in Section 4.9 is not a distribution-shift
743 artifact.

744 4.11. Cost-aware fine-tuning

745 If neither accuracy on DET variables nor the training distribution explains
746 what the supervised guidance leaves on the table, the natural suspect is the
747 objective: the rule branches on the variable it is surest about, and being sure
748 is not the same as being cheap. Three measurements test this, and the third
749 turns it into a method.

750 *R0: does the confident choice cost the least?* On 100 held-out states with
751 45–50 free variables we forced the root decision to each of the five variables
752 the network is most confident about, ran the supervised search below it,
753 and recorded the exact subtree cost of each. The confidence rule’s choice
754 is the cheapest of the five on 22% of states; the median ratio of its cost to
755 the cheapest candidate is 1.29 and the mean 2.97; on 24% of states some
756 candidate is at least twice as cheap; and choosing the cheapest candidate at
757 the root alone would remove 55% of all nodes summed over these states. The
758 LP rule’s root choice costs the same as the network’s on only 10% of states,
759 so the two rules are not merely tied. There is a large residual, and it is about
760 *which* variable to branch on.

761 *R1: can the cost be predicted?* The critic is the baseline of the policy update,
762 so it must rank decisions by cost on states it has not seen. Recording every
763 decision of the supervised search on the training pool gave 37,275 decisions
764 (9,375 on held-out states) with their exact subtree costs, of which two thirds
765 occur at 25–34 free variables. A value head on the frozen trunk reaches a
766 held-out Spearman correlation of 0.80 with $\log c$; a baseline that knows only
767 the state’s size (a linear fit on $\log |K|$) reaches 0.70, and within size buckets
768 the critic’s advantage is larger (0.68 against 0.38 at 35–44 free variables,
769 0.57 against 0.37 at 45–55). Unfreezing the trunk adds 0.02. The gate of
770 Section 3.8 is passed with the frozen critic.

771 *R2: fine-tuning.* Six epochs of 400 rollouts each ($\tau = 1$, 90 s budget, about
772 21,000 decisions per epoch) from the 45–60 free-variable training states. On
773 the 150 held-out monitor states the total subtree cost under the greedy policy,
774 relative to the supervised rule, was 0.79, 0.49, 0.56, 0.46, 0.65 and 0.56 after
775 epochs 1–6; epoch 4 was selected. Policy entropy fell from 1.73 to 0.10 over
776 the same epochs, so the policy is near-deterministic by the selected epoch.

777 *Evaluation.* Table 9 reports the fine-tuned policy (MODEL+FT) beside the
778 supervised one on the 100-instance set: 5,495 nodes against 13,386 ($2.3\times$
779 fewer, fewer on 70/100, $p < 0.001$ in every seed) and a median of 7.6 s against
780 16.5 s ($2.0\text{--}2.2\times$, faster on 64–67/100, $p \leq 0.007$), at an unchanged per-node
781 cost (738 against 796 nodes/s). Against LP-guided branching this is $4.8\text{--}5.9\times$
782 in time and $14\times$ in nodes. On the original thirty instances the effect is small
783 (8,873 against 8,657 nodes; $1.2\text{--}1.5\times$ in time, 18–21/30, p between 0.04 and
784 0.36): the seventy added instances account for the gain ($2.5\times$ fewer nodes,
785 53/70), which is consistent with the thirty being a favourable sample for the

786 supervised rule. In the deployed cascade (Table 14) the fine-tuned policy
 787 closes the lattice residual in 7.3–11.7s against 12.2–17.7s for the supervised
 788 one.

789 *Replication and transfer of the fine-tuned policy.* A second fine-tuning run
 790 with a different training seed (different rollout sampling and initialization
 791 of the correction head; epoch 2 selected by the same monitor) gives the
 792 same picture on the 100-instance set — 4,940 nodes ($2.4\times$ fewer than the
 793 supervised policy, fewer on 68/100, $p < 0.001$) and $1.9\text{--}2.4\times$ faster (66–
 794 68/100, $p \leq 0.002$) — and a larger one on the original thirty (4,040 nodes,
 795 $1.9\times$ fewer on 19/30; $2.3\times$ faster, 20–21/30). The two runs select different
 796 variables at many nodes and reach similar costs, which is what one expects
 797 if the objective, not a particular solution, is what was learned. The policy
 798 was fine-tuned on 21×60 states only, yet at 18×50 it expands 424 nodes
 799 against the supervised policy’s 1,085 ($1.7\times$ fewer, 20/30, $p = 0.099$) and is
 800 $1.8\text{--}2.2\times$ faster (22–24/30, $p \leq 0.016$): the cost-aware improvement transfers
 801 downward in size as the supervised one does.

802 *What changed.* On 300 held-out states the fine-tuned policy branches on a
 803 different root variable than the confidence rule at 95% of them. Its pick is
 804 typically the 15th most confident variable ($|\hat{p} - \frac{1}{2}|$ of 0.26 against 0.43 for the
 805 rule’s pick), and it is DET with a correct first value less often (57.5% against
 806 82.5%). The policy has moved *away* from the variables it is surest about,
 807 and the search got cheaper. The improvement is not attributable to the root
 808 decision alone: forcing only the root to the fine-tuned pick and running the
 809 supervised policy below it gives a median cost ratio of 0.93 (cheaper on 22/46
 810 disagreement states), so the gain is a property of the sequence of decisions
 811 rather than of any single one. We can therefore say what the fine-tuning
 812 optimizes — subtree cost, with the probability of being right held as a soft
 813 prior — but not yet name the feature it found; a cheap proxy, the amount of
 814 propagation the first child triggers, does not differ between the two choices.

815 4.12. Deployment in the full pipeline

816 Table 8 isolates branching by disabling lattice reduction. In deployment
 817 the lattice stage runs first and the search sees only its residual. Because
 818 that stage is independent of the branching heuristic, every arm receives the
 819 identical residual set.

820 Lattice coverage falls with size — 30/30, 24–28/30, 10–16/30 — and that
 821 decline is what creates room for branching quality to matter. At 10×25

Size	strategy	closed by lattice	residual solved by search	median time on residual	end-to-end
10×25	both	30/30	—	—	100%
18×50	LP / MODEL	24–28/30	11/11 / 11/11	0.5–5.0 s / 0.6–2.5 s	100% / 100%
21×60	RANDOM	10–16/30	0/50	> 600 s	$44.4 \pm 8.3\%$
21×60	LP	10–16/30	50/50	52.9–65.1 s	$100 \pm 0.0\%$
21×60	MODEL	10–16/30	50/50	14.2–20.2 s	$100 \pm 0.0\%$
21×60	MODEL+FT	10–16/30	50/50	7.3–11.7 s	$100 \pm 0.0\%$

Table 14: Cascade results on the warm-started loop, three seeds at every size, 600s budget at 21×60 . Lattice coverage varies across seeds (16, 14, 10 of 30 at 21×60) because the lattice stage draws random column permutations, so the residual handed to the search varies with it; all arms within a seed receive the identical residual.

and 18×50 the residual is empty or closed within seconds by either sound strategy. At 21×60 between ten and twenty instances reach the search depending on the seed, and the residual is genuinely hard: random branching closes none of the 50 residual instances pooled over seeds, so the end-to-end rate of that arm is the lattice coverage alone. Both relaxation-based and learned branching close all 50, and the difference between them is time. On the residual the learned guidance is faster on 11/14, 14/16 and 15/20 instances in the three seeds (median per-instance ratio $3.48\times$, $2.99\times$, $2.38\times$; sign test $p = 0.057$, 0.004 , 0.041), 40/50 pooled with a median ratio of $3.00\times$ ($p = 2.4 \times 10^{-5}$), expanding 6.9–9.5 $\times$ fewer nodes. The residual is where the branching heuristic operates in deployment, and the speedup there matches the isolated measurement of Table 8.

4.13. Does the graph structure contribute?

MarginalNet consumes six scalars per node and propagates them over the bipartite graph. With a feature set that small, a per-variable model ignoring the graph is a serious baseline: if it matches, the message passing is decoration. Table 15 compares models trained on identical targets, loss and splits, differing only in whether structure is used. The aggregated variant appends mean and max of the constraint-side features over each variable’s incident rows, which is the information one message-passing round would deliver.

Message passing contributes: +4.7 points over the best structure-free model and a 36% reduction in L_1 , closing roughly 60% of the residual gap to the ceiling that the MLP leaves open. The improvement in L_1 is the more relevant of the two, since the analysis of Section 3.4 consumes the marginal rather than its argmax.

Model	accuracy	L_1 to marginal	parameters
Linear, variable features only	78.1%	0.2530	4
MLP, variable features only	77.8%	0.2522	8,641
MLP + constraint aggregates	79.4%	0.2361	9,025
MarginalNet (bipartite graph)	84.1%	0.1514	21,761
Bayes ceiling	87.3%	0.0000	—

Table 15: Ablation on in-tree states of 10×25 instances, 2,000 training and 800 test states.

848 The same table bounds the claim, though. A four-parameter linear model
849 on the relaxation value alone reaches 78.1%, within 9.2 points of the ceiling;
850 most of the available signal is in that single feature, and the graph adds a
851 useful but not dominant increment on top of it.

852 4.14. Limitations

853 *Label generation bounds the method, not only its evaluation..* Exact marginals
854 require enumerating $\mathcal{S}$, which does not complete beyond $n \approx 60$ on this fam-
855 ily. Transfer permits deployment at larger sizes but not training there; ap-
856 proximating p_j by solution sampling is the natural extension and is untested.

857 *The useful size window is bounded on both sides..* Below it, lattice reduction
858 closes everything and branching quality is irrelevant. Above it ($n = 100$),
859 search fails outright for every strategy, since a constraint solver finds no
860 solution within 20 s. Our demonstration lives at $n = 60$ and we do not show
861 the window extends further.

862 *The advantage is in time and tree size, not in solve rate..* On a warm-started
863 loop both sound strategies solve every instance we evaluate within budget;
864 what the supervised guidance buys at 21×60 is a $2.0\text{--}2.3\times$ median speedup
865 and $5.5\times$ fewer nodes on 100 instances ($2.6\text{--}3.0\times$ and $9.5\times$ on the original
866 thirty, a favourable sample), and the fine-tuned policy a further $2.0\text{--}2.2\times$
867 and $2.3\times$, significant per instance in every seed. An earlier draft reported
868 a 5–0 solve-rate advantage from a loop that rebuilt the relaxation at every
869 node; that difference vanished with the faster backend and we do not claim it.
870 Whether a solve-rate advantage exists at larger budgets or sizes is untested,
871 and with 30 evaluation instances per size the power to detect one would be
872 low in any case.

873 *Multiplicity control is approximate..* Median $|\mathcal{S}|$ is 23, 19 and 10 across the
874 three sizes; at $n = 60$ multiplicity steps $42 \rightarrow 12 \rightarrow 2$ across consecutive inte-
875 ger m , so finer matching is unavailable. Three of thirty evaluation instances
876 have unverified $|\mathcal{S}|$ and are marked as such. Matching $|\mathcal{S}|$ also moves m/n
877 from 0.40 to 0.35.

878 *Cost advantage depends on scale and on baseline engineering..* The amortiza-
879 tion argument holds asymptotically but is worth little at the sizes where our
880 search operates: $1.3\text{--}1.7\times$ at the depths where prediction matters, against a
881 warm-started probing baseline that is additionally sound. It reaches $15\text{--}18\times$
882 only at $n \geq 100$, where search fails for every strategy we tested. We therefore
883 cannot demonstrate the regime in which the argument is strongest.

884 *The search implementation is not competitive..* At 21×60 CP-SAT is $9\text{--}10\times$
885 faster than the fine-tuned policy and $20\times$ faster than the supervised one on
886 the 100-instance set ($13\times$ on the original thirty). Our branching produces
887 fewer nodes, so the deficit is per-node throughput ($32\times$), now dominated by
888 the network’s forward pass; we have not shown that a competitive implemen-
889 tation is achievable: a compiled inference path, clause learning and restarts
890 are all absent, and the latter two are load-bearing for CP-SAT at larger sizes.

891 *Training and deployment distributions barely overlap, and aligning them did*
892 *not help..* The training range $\{0, \dots, 12\}$ produces states with 13–25 free
893 variables, whereas the 21×60 search queries states with 5–60 and spends
894 80% of its queries above 25. Section 4.10 shows that training on the deploy-
895 ment regime makes the search slower; we do not know whether a different
896 construction of large training states (prefixes from the search’s own trajec-
897 tory rather than from a solution, or states with several surviving solutions)
898 would behave differently.

899 *The fine-tuning result rests on one training pool and few training seeds..* The
900 policy was fine-tuned on states from one 21×60 pool; two training seeds agree
901 at 21×60 and the gain transfers to 18×50 , but no larger size was tested
902 and both seeds select their epoch on the same monitor. Its entropy collapses
903 within four epochs, epoch selection is by a held-out monitor rather than
904 by any principled stopping rule, and we cannot yet name the feature the
905 policy exploits (Section 4.11). Whether the gain transfers across sizes as the
906 supervised one does is untested.

907 *Hyperparameters are untuned and there is no validation split for the super-*
 908 *vised network..* See Section 4.2. The configuration is defensible as a first
 909 attempt and is not contaminated by the test set; the depth range, at least,
 910 was subsequently shown not to matter within a size (Section 4.10). The fine-
 911 tuning stage does select an epoch on held-out states, which are drawn from
 912 the training pool and disjoint from every evaluation instance.

913 *Single instance family..* All results concern one generator from one appli-
 914 cation. Whether the amortization argument transfers to other constraint
 915 families is untested.

916 5. Conclusion

917 We presented a learned branching heuristic for binary linear systems to-
 918 gether with an exact procedure for supervising it. Two elements carry the
 919 method. Conditioning a BLS on a partial assignment is identical to reducing
 920 it, so a single network trained on ordinary instances predicts at every node of
 921 the search tree without architectural change. And on instances small enough
 922 to enumerate, the exact conditional posterior both supplies the training tar-
 923 get and identifies which predictions can affect the search at all: only variables
 924 that every surviving solution agrees on can cost a backtrack, and their share
 925 rises from 6.3% at the root to 87.9% at depth 8.

926 Measured inside our search loop the heuristic does what it is designed
 927 to do. It is $2.0\text{--}2.3\times$ faster in median than relaxation-based branching at
 928 21×60 on 100 instances in each of three seeds, expands $5.5\times$ fewer nodes,
 929 and closes the residual of the deployed pipeline $3.0\times$ faster, where random
 930 branching closes nothing; with solution multiplicity controlled, a network
 931 trained at $2.4\times$ smaller scale performs indistinguishably from one trained at
 932 the evaluation size, which matters because exact supervision is available only
 933 at small sizes.

934 We are equally explicit about what the results do not establish. The sys-
 935 tem does not compete with a complete solver: CP-SAT is $9\text{--}40\times$ faster at
 936 every size we tested. Decomposing that gap is informative rather than excul-
 937 patory — our branching yields $2.2\times$ fewer nodes and loses on throughput by
 938 $32\times$, which places the deficit in a Python research implementation whose per-
 939 node cost is now the network’s forward pass, and which lacks clause learning
 940 and restarts — but we have not shown that closing it is achievable. Nor does
 941 the cost argument carry the weight we initially assigned it. Against a prop-
 942 erly warm-started probing baseline the advantage is $1.3\text{--}1.7\times$ at the depths

943 where prediction matters on our instances, and probing is sound where the
 944 network is not; the ratio reaches 15–18 $\times$ only at $n \geq 100$, a regime in which
 945 search fails for every strategy we tried. An earlier version of this compar-
 946 ison, using a backend that rebuilt the LP for each probe and excluded the
 947 network’s own feature solve from its timer, reported 19–23 $\times$ at a scale where
 948 the true figure is near unity, and we record the correction because the error
 949 came entirely from instrumentation rather than from what the two methods
 950 compute.

951 Two further results sharpen what the supervised target is. Moving the
 952 training states to the deployment regime, where the network’s prediction
 953 margin lives, does not help and in fact doubles the search tree, because the
 954 search spends its nodes in the small interior states the moved network no
 955 longer handles well; the training-state distribution is not the lever. Optimiz-
 956 ing the cost the search pays is. With the exact subtree cost of every decision
 957 read off the search, an actor-critic fine-tuning of the branching choice —
 958 initialized at the supervised rule, with the marginal network frozen — ex-
 959 pands 2.3 $\times$ fewer nodes and runs 2.0–2.2 $\times$ faster than the supervised policy
 960 on 100 instances, and the policy it finds branches on variables the network
 961 is *less* sure about. The probability that a branch is right and the cost of
 962 the branch being wrong are different objectives, and the second is the one a
 963 search minimizes.

964 What remains, and what we claim, is methodological. For planted-
 965 solution instance families the exact conditional posterior is computable and is
 966 a well-posed supervision signal for branching, where the conventional target
 967 — a single planted solution — is not. States drawn from inside the search
 968 tree are a better training distribution than root states, and the reduction
 969 identity makes generating them a matter of substitution rather than solver
 970 instrumentation. The resulting heuristic transfers across instance size when
 971 solution multiplicity is held comparable, and the exact cost of its decisions,
 972 available from the search, is a second well-posed training signal that im-
 973 proves it. Whether this becomes practically useful depends on two things
 974 we have not done: extending exact supervision past $n \approx 60$, most plausibly
 975 by approximating the marginals without exhaustive enumeration, and imple-
 976 menting the search well enough that a branching improvement is visible at
 977 the system level.

978 Reproducibility

979 Code is organized by component under `proposed_src/`, with shared mod-
980 ules in `util/` and a `ROLES.txt` in each directory describing every file. Scripts
981 are run from their own directory. Every search number in Tables 8, 12
982 and 14 comes from one chain, `v23_ablation/run_revision2.sh`, which runs
983 the warm-started loop (`v23_warm_search.py`) over all conditions and seeds;
984 `v23_aggregate.py` produces the tables from its output.

```
985 # instances with matched |S| and exact marginals
986 cd proposed_src/v15_bayes
987 python3 v15_gen_by_m.py --m 10 --n 25 --count 2200 --K 20 \
988     --out ../../runs/v22/sols_10x25.json
989
990 # attach solution sets, then split 80/20 (seed 0)
991 cd ../v22_transfer
992 python3 v22_add_solutions.py \
993     --inp ../../runs/v22/sols_10x25.json \
994     --out ../../runs/v22/full_10x25.json
995 python3 v22_split.py # -> full_10x25_{train,test}
996
997 # in-tree training states (depths 0-12, 6 per instance)
998 cd ../v17_conditional
999 python3 v17_make_conditional_data.py \
1000     --sols ../../runs/v22/full_10x25_train.json \
1001     --max_depth 12 --per_instance 6 \
1002     --out ../../runs/v22/cond_10x25_train.json
1003
1004 # train (about 20 minutes, single CPU core)
1005 python3 v17_train_conditional.py \
1006     --train ../../runs/v22/cond_10x25_train.json \
1007     --test  ../../runs/v22/cond_10x25_test.json \
1008     --n_train 10000 --n_test 2000 --epochs 20 \
1009     --target soft --out ../../runs/model_n25
1010
1011 # every search table: warm-started loop, 3 seeds,
1012 # conditions A/B/R/C/D/Crand x {lattice off, on}
1013 cd ../v23_ablation
1014 bash run_revision2.sh # -> runs/v23/, chain.log
1015 python3 v23_aggregate.py # -> runs/v23/aggregate.json
```

```

1016
1017 # v24: range training (negative), 100-instance
1018 # set, R0/R1/R2 fine-tuning and evaluation
1019 cd ../v24_deploy_range
1020 bash run_v24.sh; bash run_v24_switch.sh
1021 bash run_v24_r1.sh; bash run_v24_r2.sh
1022 python3 v24_aggregate.py \
1023     --tags M0 M1 M2 SW PV

```

1024 All generators are seeded. The released checkpoints are 96 KB each. In-
1025 stance data and weights are regenerated from the above rather than dis-
1026 tributed, and the full evaluation cycle is scripted in `run_revision2.sh` under
1027 `v23_ablation/`.

1028 References

- 1029 [1] Aardal, K., Hurkens, C. A. J., and Lenstra, A. K. (2000). Solving a
1030 system of linear Diophantine equations with lower and upper bounds on
1031 the variables. *Mathematics of Operations Research*, 25(3):427–442.
- 1032 [2] Aardal, K., Bixby, R. E., Hurkens, C. A. J., Lenstra, A. K., and
1033 Smeltink, J. W. (2000). Market split and basis reduction: Towards a
1034 solution of the Cornuéjols–Dawande instances. *INFORMS Journal on*
1035 *Computing*, 12(3):192–202.
- 1036 [3] Amos, B. (2023). Tutorial on amortized optimization. *Foundations and*
1037 *Trends in Machine Learning*, 16(5):592–732.
- 1038 [4] Bengio, Y., Lodi, A., and Prouvost, A. (2021). Machine learning for
1039 combinatorial optimization: A methodological tour d’horizon. *European*
1040 *Journal of Operational Research*, 290(2):405–421.
- 1041 [5] Chen, Z., Liu, J., Wang, X., Lu, J., and Yin, W. (2023). On represent-
1042 ing mixed-integer linear programs by graph neural networks. In *Inter-*
1043 *national Conference on Learning Representations*.
- 1044 [6] Cornuéjols, G. and Dawande, M. (1999). A class of hard small 0-1 pro-
1045 grams. In *Integer Programming and Combinatorial Optimization*, pages
1046 284–293.

- 1047 [7] Gasse, M., Chételat, D., Ferroni, N., Charlin, L., and Lodi, A. (2019).
1048 Exact combinatorial optimization with graph convolutional neural net-
1049 works. In *Advances in Neural Information Processing Systems*.
- 1050 [8] Gershman, S. J. and Goodman, N. D. (2014). Amortized inference in
1051 probabilistic reasoning. In *Proceedings of the 36th Annual Meeting of
1052 the Cognitive Science Society*, pages 517–522.
- 1053 [9] Khalil, E. B., Le Bodic, P., Song, L., Nemhauser, G., and Dilkina, B.
1054 (2016). Learning to branch in mixed integer programming. In *AAAI
1055 Conference on Artificial Intelligence*, pages 724–731.
- 1056 [10] Lenstra, A. K., Lenstra, H. W., and Lovász, L. (1982). Factoring poly-
1057 nomials with rational coefficients. *Mathematische Annalen*, 261(4):515–
1058 534.
- 1059 [11] Li, Z., Guo, J., and Si, X. (2023). G4SATBench: Benchmarking and
1060 advancing SAT solving with graph neural networks. *Transactions on
1061 Machine Learning Research*.
- 1062 [12] Marques-Silva, J. P. and Sakallah, K. A. (1999). GRASP: A search algo-
1063 rithm for propositional satisfiability. *IEEE Transactions on Computers*,
1064 48(5):506–521.
- 1065 [13] Millidge, B. (2022). Deconfusing direct vs amortized opti-
1066 mization. Unrefereed technical note, [https://www.beren.io/
1067 2022-09-25-Deconfusing-direct-vs-amortized-optimization/](https://www.beren.io/2022-09-25-Deconfusing-direct-vs-amortized-optimization/).
- 1068 [14] Moskewicz, M. W., Madigan, C. F., Zhao, Y., Zhang, L., and Malik, S.
1069 (2001). Chaff: Engineering an efficient SAT solver. In *Design Automa-
1070 tion Conference*, pages 530–535.
- 1071 [15] Ohrimenko, O., Stuckey, P. J., and Codish, M. (2009). Propagation via
1072 lazy clause generation. *Constraints*, 14(3):357–391.
- 1073 [16] Schnorr, C. P. and Euchner, M. (1994). Lattice basis reduction: Im-
1074 proved practical algorithms and solving subset sum problems. *Mathe-
1075 matical Programming*, 66(1–3):181–199.
- 1076 [17] Selsam, D., Lamm, M., Bünz, B., Liang, P., de Moura, L., and Dill,
1077 D. L. (2019). Learning a SAT solver from single-bit supervision. In *In-
1078 ternational Conference on Learning Representations*.

- 1079 [18] Veličković, P., Cucurull, G., Casanova, A., Romero, A., Liò, P., and Ben-
1080 gio, Y. (2018). Graph attention networks. In *International Conference*
1081 *on Learning Representations*.
- 1082 [19] Williams, R. J. (1992). Simple statistical gradient-following algorithms
1083 for connectionist reinforcement learning. *Machine Learning*, 8(3–4):229–
1084 256.
- 1085 [20] Etheve, M., Alès, Z., Bissuel, C., Juan, O., and Kedad-Sidhoum, S.
1086 (2020). Reinforcement learning for variable selection in a branch and
1087 bound algorithm. In *Integration of Constraint Programming, Artificial*
1088 *Intelligence, and Operations Research (CPAIOR)*, pages 176–185.
- 1089 [21] Scavuzzo, L., Chen, F. Y., Chételat, D., Gasse, M., Lodi, A., Yorke-
1090 Smith, N., and Aardal, K. (2022). Learning to branch with tree MDPs.
1091 In *Advances in Neural Information Processing Systems*, volume 35.
- 1092 [22] Haralick, R. M. and Elliott, G. L. (1980). Increasing tree search efficiency
1093 for constraint satisfaction problems. *Artificial Intelligence*, 14(3):263–
1094 313.
- 1095 [23] Cappart, Q., Moisan, T., Rousseau, L.-M., Prémont-Schwarz, I., and
1096 Cire, A. A. (2021). Combining reinforcement learning and constraint
1097 programming for combinatorial optimization. In *AAAI Conference on*
1098 *Artificial Intelligence*.
- 1099 [24] Chu, G. and Stuckey, P. J. (2015). Learning value heuristics for con-
1100 straint programming. In *CPAIOR*.
- 1101 [25] Feng, S. and Yang, Y. (2025). SORREL: Suboptimal-demonstration-
1102 guided reinforcement learning for learning to branch. In *AAAI Confer-*
1103 *ence on Artificial Intelligence*.
- 1104 [26] Han, Q., Yang, L., Chen, Q., Zhou, X., Zhang, D., Wang, A., Sun,
1105 R., and Luo, X. (2023). A GNN-guided predict-and-search framework
1106 for mixed-integer linear programming. In *International Conference on*
1107 *Learning Representations*.
- 1108 [27] Khalil, E. B., Morris, C., and Lodi, A. (2022). MIP-GNN: A data-driven
1109 framework for guiding combinatorial solvers. In *AAAI Conference on*
1110 *Artificial Intelligence*.

- 1111 [28] Li, Z. and Si, X. (2022). NSNet: A general neural probabilistic frame-
1112 work for satisfiability problems. In *Advances in Neural Information Pro-*
1113 *cessing Systems*, volume 35.
- 1114 [29] Montanari, A., Ricci-Tersenghi, F., and Semerjian, G. (2007). Solving
1115 constraint satisfaction problems through belief propagation-guided dec-
1116 imation. In *Allerton Conference on Communication, Control, and Com-*
1117 *puting*.
- 1118 [30] Nair, V., Bartunov, S., Gimeno, F., et al. (2020). Solving mixed integer
1119 programs using neural networks. *arXiv:2012.13349*.
- 1120 [31] Parsonson, C. W. F., Laterre, A., and Barrett, T. D. (2023). Reinforce-
1121 ment learning for branch-and-bound optimisation using retrospective
1122 trajectories. In *AAAI Conference on Artificial Intelligence*.
- 1123 [32] Pesant, G., Quimper, C.-G., and Zanarini, A. (2012). Counting-based
1124 search: Branching heuristics for constraint satisfaction problems. *Jour-*
1125 *nal of Artificial Intelligence Research*, 43:173–210.
- 1126 [33] Refalo, P. (2004). Impact-based search strategies for constraint pro-
1127 gramming. In *Principles and Practice of Constraint Programming (CP)*,
1128 pages 557–571.
- 1129 [34] Vaezipoor, P., Lederman, G., Wu, Y., Maddison, C., Grosse, R. B., Se-
1130 shia, S. A., and Bacchus, F. (2021). Learning branching heuristics for
1131 propositional model counting. In *AAAI Conference on Artificial Intelli-*
1132 *gence*.
- 1133 [35] Wassermann, A. (2025). Solving the market split problem with lattice
1134 enumeration. *arXiv:2508.08702*; to appear in *Mathematical Program-*
1135 *ming Computation*.
- 1136 [36] Xu, K., Hu, W., Leskovec, J., and Jegelka, S. (2019). How powerful
1137 are graph neural networks? In *International Conference on Learning*
1138 *Representations*.