

Amortized Deduction for Binary Linear Systems: Learning Branching Guidance from Exact Conditional Marginals

neuroMCTS project^a

^a*Korea University, Seoul, Republic of Korea*

Abstract

A binary linear system (BLS) asks for $\mathbf{x} \in \{0, 1\}^n$ satisfying $A\mathbf{x} = \mathbf{b}$ with $A \in \{0, 1\}^{m \times n}$, a formulation that arises in laser-based non-destructive testing, where a binary occupancy image is reconstructed from integer projections. On market-split-style instances the linear relaxation is deliberately uninformative, so backtracking search spends most of its effort on branching decisions that a cheap heuristic cannot make well. We present a learned branching heuristic for this setting and an exact procedure for training it.

The method rests on two observations. First, conditioning a BLS on a partial assignment is *identical* to substituting the assignment out and obtaining a smaller BLS; a single network therefore predicts at the root and at every interior node without architectural change, and training states are ordinary instances of reduced size. Second, on instances small enough to enumerate, the exact conditional marginal $p_j = \Pr[x_j = 1 \mid A, \mathbf{b}]$ is computable, and it partitions the free variables of a node into those every surviving solution agrees on — where branching the other way empties the subtree and forces a backtrack — and those solutions disagree on, where either branch still reaches a solution. Only the first kind carries cost. This partition supplies both an exact training target and a criterion for which predictions matter: at the root only 6.3% of variables are of the first kind, but by depth 8 the share is 87.9%, which is where prediction accuracy begins converting into search cost.

We train a 21.8K-parameter bipartite graph network on these targets and use it for branching only, never for fixing variables, since it produces a probability rather than a proof. Against LP-probing — the sound procedure that detects the same variables, at one linear program each — the network needs

a single relaxation and one forward pass, so their cost ratio grows linearly in instance size: $1.7\times$ at $n = 25$, $7.2\times$ at $n = 60$, $17.8\times$ at $n = 150$, measured with a warm-started backend on both sides. Under a matched wall-clock budget on a warm-started search loop, replicated over three seeds, the resulting search is $2.6\text{--}3.0\times$ faster in median than LP-guided branching at 21×60 (faster on 23/30 instances in every seed, sign test $p = 0.005$) and expands $9.5\times$ fewer nodes, with both strategies solving every instance; within the full pipeline, where lattice reduction runs first, it closes the residual $3.0\times$ faster than LP-guided branching, while random branching closes none of it. Controlling the number of solutions across sizes, we find no measurable transfer penalty: a network trained only at 10×25 is statistically indistinguishable from one trained at the evaluation size. On a 100-instance 21×60 set the supervised guidance is $2.0\text{--}2.3\times$ faster than LP-guided branching with $5.5\times$ fewer nodes. Two further experiments locate what the supervised target does and does not capture. Training on states of deployment size — the regime where the network’s prediction margin lives — does *not* help: it loses competence on the small interior states where most nodes are spent, and the search gets slower. What does help is optimizing the objective the search actually pays: with the exact subtree cost of every branching decision available from the search itself, an actor-critic fine-tuning of the branching policy, initialized at the supervised rule, expands $2.3\times$ fewer nodes and is $2.0\text{--}2.2\times$ faster than the supervised policy on the same 100 instances ($p < 0.01$), $4.8\text{--}5.9\times$ faster than LP-guided branching. The fine-tuned policy branches on a different variable than the confidence rule at 95% of held-out states, typically one it is *less* sure about, which is direct evidence that the probability of being right and the cost of being wrong are different objectives.

We also report what the method does not achieve. CP-SAT solves the same instances $9\text{--}40\times$ faster at every size tested ($9\text{--}10\times$ at 21×60 after fine-tuning), and our system does not compete with it; decomposing that gap shows our branching produces $2.2\times$ *fewer* nodes but processes them $32\times$ more slowly, which places the deficit in per-node cost rather than in branching quality, though we do not show that it is recoverable. And the cost advantage over probing, while asymptotically real, is only $1.3\text{--}1.7\times$ at the depths where prediction matters most on our instances, where probing additionally offers proofs rather than predictions. The contribution we claim is therefore methodological: exact conditional posteriors are an obtainable and well-posed supervision signal for branching, in-tree states are a better training distribution than root states — a root-trained network is worse than

relaxation rounding when deployed in the tree, an in-tree-trained one is better by 3.4 points — and the resulting heuristic transfers across instance size. Code, generators, exact-marginal labels and per-instance results are released.

Keywords: Binary linear systems, Learning to branch, Amortized optimization, Exact conditional marginals, Market split, Graph neural networks

1. Introduction

1.1. Problem

Let $A \in \{0, 1\}^{m \times n}$ and $\mathbf{b} \in \mathbb{Z}^m$. A *binary linear system* (BLS) is the feasibility problem

$$\text{find } \mathbf{x} \in \{0, 1\}^n \text{ such that } A\mathbf{x} = \mathbf{b}, \quad (1)$$

whose solution set we write $\mathcal{S} = \{\mathbf{x} \in \{0, 1\}^n : A\mathbf{x} = \mathbf{b}\}$. Each row is a cardinality constraint $\sum_j a_{ij}x_j = b_i$ stating how many cells of a designated subset are occupied.

The motivating application is laser-based non-destructive testing. A specimen is scanned along m directions; each scan returns the integer count of absorbing cells along its path, and the task is to reconstruct the binary occupancy image. Two properties of that setting shape the problem formulation. Solutions need not be unique — several images can be consistent with the same measurements, and recovering *any* member of $\mathcal{S}$ solves the physical problem. And measurement noise can perturb $\mathbf{b}$ so that $\mathcal{S} = \emptyset$, making a reliable infeasibility verdict as valuable as a solution.

Deciding $\mathcal{S} \neq \emptyset$ is NP-complete in general. The instances studied here are drawn from a deliberately hard regime modelled on market-split problems [6, 2], in which the linear relaxation

$$\tilde{\mathcal{S}} = \{\mathbf{x} \in [0, 1]^n : A\mathbf{x} = \mathbf{b}\} \quad (2)$$

is a large polytope with predominantly fractional vertices. Relaxation-based bounding is therefore weak, and a complete solver must decide where to branch with little guidance from the relaxation — exactly the decision this paper learns.

1.2. Where a learned component attaches, and what to train it on

A complete solver for (1) alternates deduction (constraint propagation, relaxation-based pruning) with a *branching decision*: which free variable to split on, and which value to try first. Deduction is sound and must remain so; branching is a heuristic choice, where an error costs search time but never correctness. This asymmetry makes branching the natural place to attach a learned component, and it is where learning has been most successful in mixed-integer programming [9, 7].

The difficulty is supervision. Two standard answers are available and both are unsatisfactory here. Training against a planted solution $\mathbf{x}^*$ treats one member of $\mathcal{S}$ as ground truth, which is ill-posed when $|\mathcal{S}| > 1$: nothing in $(A, \mathbf{b})$ distinguishes the planted solution from any other, so the label is a sample from the posterior rather than the posterior. Imitating an expensive expert, as in strong-branching imitation, requires an expert whose decisions are both informative and affordable to collect; the natural candidate here, LP-probing, costs one linear program per variable per node.

We take a third route, available because these instances have planted solutions and because feasibility can be enumerated at moderate size: compute the *exact conditional posterior* by exhaustive enumeration on small instances, and train against it. Section 3 shows that this is both tractable and, in a precise sense, the right target — and that the posterior additionally identifies *which* predictions can affect search at all.

1.3. Contributions

1. **A reduction identity that makes in-tree supervision trivial to implement.** Conditioning (1) on a partial assignment yields a smaller BLS with the same structure (Proposition 1). One architecture, one label generator and one training loop therefore serve the root and every interior node, and a depth- d training state is stored as an ordinary instance on $n - d$ variables.
2. **A partition of branching decisions by whether they can cost anything.** The exact marginals split a node’s free variables into DET (all surviving solutions agree; branching the other way empties the subtree) and OPEN (solutions disagree; either branch reaches a solution). Only DET errors cost a backtrack. We measure the composition: 6.3% DET at the root, rising to 87.9% at depth 8 (Sections 4.3, 4.4). This locates useful prediction inside the tree rather than at the root and explains why root-level accuracy is a poor objective.

- 60 3. **A cost characterization of the learned heuristic against the**
 61 **sound procedure it approximates.** The network needs one relax-
 62 ation and one forward pass; LP-probing needs one relaxation per vari-
 63 able. Measured with a warm-started backend on both sides, the ratio
 64 scales from $1.7\times$ at $n = 25$ to $17.8\times$ at $n = 150$ (Section 4.5). We
 65 report this as a scaling law rather than a single figure, and note that
 66 at the sizes where our search operates the advantage is modest and
 67 probing has the additional merit of soundness.
- 68 4. **End-to-end evaluation under matched wall-clock budgets.** At
 69 21×60 on a warm-started search loop, replicated over three seeds, the
 70 learned guidance is $2.0\text{--}2.3\times$ faster in median than LP-based branching
 71 on a 100-instance set (faster on 70–71/100, sign test $p < 0.001$) and
 72 expands $5.5\times$ fewer nodes; on the original 30-instance set the margins
 73 are $2.6\text{--}3.0\times$ and $9.5\times$, and both strategies solve every instance within
 74 600 s (Section 4.6). In the deployed cascade both sound strategies close
 75 every residual instance, the learned one $3.0\times$ faster (40/50 wins pooled
 76 over seeds, $p = 2 \times 10^{-5}$), while random branching closes none of the
 77 50 residual instances (Section 4.12).
- 78 5. **Size transfer with solution multiplicity controlled.** Holding $|\mathcal{S}|$
 79 comparable across sizes, a network trained only at 10×25 matches one
 80 trained at the evaluation size 21×60 (faster on 15/30 instances; total
 81 time differing by 0.4%), showing the learned quantity is not size-specific
 82 (Section 4.8).
- 83 6. **A negative result on aligning the training distribution with**
 84 **deployment.** The network’s margin over relaxation rounding lives on
 85 states with 50–60 free variables, outside its 13–25 training range. Train-
 86 ing on exactly those states, with everything else fixed, raises accuracy
 87 there by about one point and *doubles* the search tree, because it loses
 88 competence on the small interior states where two thirds of all nodes
 89 are spent (Section 4.10).
- 90 7. **Cost-aware fine-tuning of the branching policy.** The exact cost
 91 of a branching decision — the number of nodes its subtree expands —
 92 is available from the search itself, and on held-out states the confidence
 93 rule picks the cheapest of its five candidates only 22% of the time.
 94 An actor-critic fine-tuning on that cost, initialized at the supervised
 95 rule with the marginal network frozen, expands $2.3\times$ fewer nodes and
 96 is $2.0\text{--}2.2\times$ faster than the supervised policy on the 100-instance set
 97 (70/100 and 64–67/100 wins, $p < 0.01$ in every seed), $4.8\text{--}5.9\times$ faster

98 than LP-guided branching (Section 4.11).

99 2. Related Work

100 2.1. Binary linear systems and market-split hardness

101 Problem (1) is a 0/1 integer feasibility problem and is NP-complete for
102 general A . Practical hardness, however, depends on the instance distribution.
103 Cornuéjols and Dawande [6] introduced the market-split family, in which m
104 agents must divide n items so that every agent’s budget is met exactly; with
105 $n \approx 10(m - 1)$ these instances defeat branch-and-bound at remarkably small
106 sizes. The mechanism is the integrality gap: the relaxation polytope is large
107 and its vertices are fractional, so bounding prunes little and the search must
108 branch nearly blindly. Aardal et al. [2] showed that lattice basis reduction
109 attacks the same family far more effectively than branch-and-bound, a result
110 we rely on in Section 3.2.

111 Our generator (Section 4.1) makes this mechanism explicit by maximizing
112 the spread of relaxation vertices, which enlarges $\tilde{\mathcal{S}}$ and weakens the relaxation
113 signal by construction.

114 2.2. Complete solvers

115 *Mixed-integer programming.* Branch-and-bound over the linear relaxation,
116 strengthened by presolve and cutting planes, is the standard approach. Its
117 leverage comes from pruning nodes whose relaxation is infeasible or bounded
118 away from the target, which is precisely what the market-split construction
119 removes.

120 *Constraint programming with clause learning.* CP-SAT-style solvers com-
121 bine constraint propagation with conflict-driven clause learning [12, 14] and
122 lazy clause generation [15]. Their strength on this family comes from a
123 property central to our argument: a search *manufactures* information as it
124 descends. A marginal $\Pr[x_j = 1 \mid A, \mathbf{b}]$ may be uninformative at the root
125 while $\Pr[x_j = 1 \mid A, \mathbf{b}, x_1=1, x_5=0]$ is 0 or 1. Structure invisible at the root
126 becomes deducible after commitments are made, which is why Section 3.4
127 studies the interior of the tree rather than the root.

128 *Lattice reduction..* An alternative attack embeds (1) in a lattice and searches
 129 for short vectors. LLL [10] reduces a basis in polynomial time and BKZ
 130 [16] generalizes it with a block size trading time against basis quality. The
 131 Aardal–Hurkens–Lenstra embedding [1, 2] is constructed so that a 0/1 so-
 132 lution appears directly as a short vector, making reduction a certificate-
 133 producing procedure rather than a heuristic. We use it as the first stage of
 134 the deployed pipeline.

135 2.3. Learning for combinatorial optimization

136 Bengio et al. [4] survey this field and frame it in terms we adopt: solvers
 137 rely on hand-crafted heuristics for decisions that are “otherwise too expensive
 138 to compute or mathematically not well defined,” and learning is a candidate
 139 for making those decisions better. The phrase *too expensive to compute*
 140 presumes the decision is computable in principle and asks only about cost,
 141 which is the situation we measure in Section 4.5.

142 *Branching imitation..* The clearest success in this area is imitation of an
 143 expensive expert. Strong branching scores candidates by tentatively solving
 144 both children’s relaxations; it yields small trees and is too slow to run at
 145 every node. Gasse et al. [7] train a bipartite graph network to imitate it and
 146 recover most of the benefit at a fraction of the cost, following Khalil et al.
 147 [9]. The gain is not new information — the expert already knew the answer
 148 — but cheap access to it. Our method occupies the same structural position,
 149 with two differences: the target is an exact posterior rather than an expert’s
 150 scores, and the analysis identifies which of the expert’s decisions matter.

151 *Direct satisfiability prediction..* Attempts to have a network decide satisfia-
 152 bility end-to-end have a weaker record. NeuroSAT [17] classifies satisfiability
 153 by message passing over a literal–clause graph and succeeds on small ran-
 154 dom instances, but Li et al. [11] report substantial accuracy collapse under
 155 size transfer. Chen et al. [5] prove that standard message-passing networks
 156 cannot separate certain feasible/infeasible MILP pairs at all, since the pairs
 157 are indistinguishable in the Weisfeiler–Lehman hierarchy that bounds such
 158 architectures [23]. We therefore do not ask the network for a feasibility ver-
 159 dict; it supplies branching guidance inside a complete search that remains
 160 responsible for soundness.

161 *Graph representations..* We use graph attention [18] over the bipartite variable–
 162 constraint graph, the standard encoding for MILP in this literature [7].

163 *Reinforcement learning for branching..* Imitation targets an expert’s ranking,
 164 not the search cost the ranking is meant to reduce. Ethève et al. [20] and
 165 Scavuzzo et al. [21] instead optimize tree size directly, the latter formalizing
 166 branch-and-bound as a tree MDP in which the return of a decision is the
 167 size of its subtree, so that credit is assigned per decision rather than per
 168 episode. Our fine-tuning stage (Section 3.8) uses the same per-subtree credit
 169 with a plain policy-gradient update [19], and differs in its starting point:
 170 the policy is initialized at a supervised rule derived from exact conditional
 171 marginals, and the marginal network is frozen, so the only thing the cost
 172 signal can change is which variable is branched on. The classical statement
 173 of the objective is the fail-first principle [22]: choose the variable whose wrong
 174 branch is refuted soonest.

175 2.4. Amortized optimization

176 The distinction organizing our results predates its use in combinatorial
 177 optimization. Gershman and Goodman [8] introduced *amortized inference*:
 178 when many related queries must be answered, solving each from scratch is
 179 wasteful and the question becomes how to reuse computation across queries.
 180 Amos [3] develops the same idea as *amortized optimization* — using learning
 181 to predict solutions by exploiting shared structure across similar instances
 182 — and surveys its appearance in variational inference, meta-learning and
 183 control. Millidge [13] states the trade-off in the form most relevant here:
 184 *direct* optimization spends computation on the instance at hand and pays
 185 that cost at every invocation, whereas *amortized* optimization converts the
 186 problem into supervised learning with cheap inference but requires a corpus
 187 of solved instances up front and is bounded by the approximator’s gener-
 188 alization; the productive arrangement is hybrid, with amortized predictions
 189 steering a direct optimizer.

190 Our system instantiates that hybrid precisely. The direct method is LP-
 191 probing, sound and costly. The amortized method is the network, cheap and
 192 unsound. The hybrid rule is that the network steers branching while probing
 193 and propagation retain exclusive authority over variable fixing (Section 3.7).

194 3. Proposed Method

195 3.1. Setting and notation

196 For a partial assignment fixing variables $F \subseteq \{1, \dots, n\}$ to values $\mathbf{v} \in$
 197 $\{0, 1\}^{|F|}$, write

$$\mathcal{S}(F, \mathbf{v}) = \{\mathbf{x} \in \mathcal{S} : \mathbf{x}_F = \mathbf{v}\} \quad (3)$$

198 for the solutions consistent with it, and let $K = \{1, \dots, n\} \setminus F$ index the free
 199 variables. The *conditional marginal* of a free variable $j \in K$ is

$$p_j(F, \mathbf{v}) = \frac{|\{\mathbf{x} \in \mathcal{S}(F, \mathbf{v}) : x_j = 1\}|}{|\mathcal{S}(F, \mathbf{v})|}. \quad (4)$$

200 3.2. System overview

201 The deployed solver is a cascade of a lattice stage and a guided complete
 202 search:

$$\text{propagation} \rightarrow \text{LP infeasibility rule} \rightarrow \text{kernel pump} \rightarrow \text{AHL/BKZ} \rightarrow \text{guided backtracking search} \quad (5)$$

203 The first three stages are cheap and sound. AHL/BKZ attempts a certificate-
 204 producing lattice reduction with a size-dependent block parameter; when
 205 it succeeds it returns a verified solution outright. The learned component
 206 appears only in the final stage, and only in its branching decisions. This
 207 placement matters for the evaluation: because the lattice stage is independent
 208 of the branching heuristic, every heuristic under comparison receives exactly
 209 the same residual set of instances, so conditioning on that residual describes
 210 the algorithm rather than selecting favourable cases.

211 3.3. Conditioning is reduction

212 **Proposition 1.** *Let A_F and A_K denote the column submatrices of A indexed*
 213 *by F and K . Then*

$$\mathcal{S}(F, \mathbf{v}) \cong \{\mathbf{y} \in \{0, 1\}^{|K|} : A_K \mathbf{y} = \mathbf{b} - A_F \mathbf{v}\}, \quad (6)$$

214 *the bijection being $\mathbf{x} \mapsto \mathbf{x}_K$. Consequently the conditional marginals (4)*
 215 *of the original instance are the ordinary marginals of the reduced instance*
 216 *$(A_K, \mathbf{b} - A_F \mathbf{v})$.*

217 *Proof.* $A\mathbf{x} = A_F \mathbf{x}_F + A_K \mathbf{x}_K$. Fixing $\mathbf{x}_F = \mathbf{v}$ gives $A\mathbf{x} = \mathbf{b} \iff A_K \mathbf{x}_K =$
 218 $\mathbf{b} - A_F \mathbf{v}$, and $\mathbf{x} \mapsto \mathbf{x}_K$ is a bijection between the two sets since $\mathbf{x}_F$ is deter-
 219 mined. The marginal statement follows by counting. $\square$

220 The identity is elementary but it is what makes in-tree supervision practical. A depth- d search node is not a special object requiring its own encoding; 221 it is an instance of the same problem on $n - d$ variables. One network, one 222 feature extractor and one label generator therefore cover the whole tree, and 223 training data at depth d is produced by substitution rather than by instrumenting a solver. 224

225 Two implementation details preserve correctness. Rows that become all-zero after substitution carry no constraint and are removed; retaining them 226 causes linear-programming backends to reject the model. And the free bound 227 condition $0 \leq \mathbf{b} - A_F \mathbf{v} \leq A_K \mathbf{1}$ is checked before any relaxation is solved, 228 since it detects a large fraction of dead nodes at negligible cost. 229

231 3.4. Which predictions can cost anything

232 At a node with surviving solution set $\mathcal{S}(F, \mathbf{v}) \neq \emptyset$, partition the free 233 variables by their conditional marginal:

$$\text{DET} = \{j \in K : p_j \in \{0, 1\}\}, \quad \text{OPEN} = \{j \in K : 0 < p_j < 1\}. \quad (7)$$

234 The distinction is operational rather than statistical:

- 235 • For $j \in \text{OPEN}$, both children of a branch on x_j contain solutions. Either 236 choice keeps a solution reachable, so *no branching decision on a OPEN 237 variable can cost anything.*
- 238 • For $j \in \text{DET}$, one child is empty. Branching against the forced value 239 guarantees that the subtree is exhausted without a solution, costing a 240 full backtrack.

241 This gives the conditional marginal its operational reading. It is not a 242 “probability of being correct” about some designated solution, but the *fraction of the surviving solution set retained* by an assignment; and the only 243 errors that matter are those on variables where that fraction is 0 or 1. 244

245 Two consequences follow, both of which the experiments confirm. First, 246 an aggregate accuracy over all free variables is a poor objective, because it 247 mixes decisions that can cost a backtrack with decisions that cannot; the 248 relevant quantity is accuracy on DET. Second, since $|\mathcal{S}(F, \mathbf{v})|$ shrinks as the 249 search descends, the DET share grows with depth. This alone does not say 250 *where* a learned predictor is useful, because a DET variable that the relaxation 251 already rounds correctly needs no predictor. The useful region is where DET

variables are plentiful *and* the relaxation is weak, and Section 4.9 shows that this is governed by constraint density $m/|K|$ rather than by depth as such: on small instances the two coincide, on larger ones they do not.

3.5. Exact conditional marginals as training targets

For instances small enough to enumerate, $\mathcal{S}$ is obtained exactly with a constraint solver in all-solutions mode, and (4) is computed by counting. Training minimizes the cross-entropy to the resulting Bernoulli targets,

$$\mathcal{L} = -\frac{1}{|K|} \sum_{j \in K} \left[p_j \log \hat{p}_j + (1 - p_j) \log(1 - \hat{p}_j) \right], \quad (8)$$

which is minimized exactly at $\hat{p}_j = p_j$. Note that the same loss with a planted solution in place of p_j — the conventional choice — is a noisy surrogate whenever $|\mathcal{S}| > 1$, since the planted solution is one draw from the posterior; on our instances a Bayes-optimal predictor agrees with it only about two thirds of the time.

Enumeration correctness.. A subtlety here is easy to get wrong and silently corrupts every label derived from it. Constraint solvers report distinct statuses for “search exhausted” and “time limit reached with solutions already collected.” Only the former certifies that $\mathcal{S}$ is complete; accepting the latter substitutes a truncated solution set, which biases every marginal computed from it. In our implementation only OPTIMAL and INFEASIBLE are accepted and timed-out instances are discarded rather than averaged in. A useful diagnostic is that mean enumeration time equal to the time limit indicates truncated instances being counted as complete.

Sampling reachable states.. A depth- d training state is built by drawing d uniformly from $\{0, \dots, 12\}$, taking a prefix of an LP-confidence variable ordering, and fixing those variables to the values of a randomly chosen *actual solution*. Fixing arbitrary values would produce states with $\mathcal{S}(F, \mathbf{v}) = \emptyset$; such states are dead nodes that propagation prunes immediately and on which the conditional marginal is undefined, so they are not states a branching heuristic is ever asked about. The depth range is a design choice we did not tune; on 10×25 instances it yields training states with 13–25 free variables, and Section 4.9 examines what happens when the deployed search queries states outside that range, which at 21×60 is most of them.

283 3.6. Network

284 An instance is encoded as a bipartite graph with n variable nodes, m
 285 constraint nodes, and an edge (j, i) whenever $a_{ij} = 1$. Let $\tilde{\mathbf{x}}$ be the relaxation
 286 solution of the current (reduced) instance and $r_i = \sum_j a_{ij}$ the row sums.
 287 Node features are

$$\text{variable } j : (\tilde{x}_j, \deg(j)/m, |\tilde{x}_j - \text{round}(\tilde{x}_j)|), \quad (9)$$

$$\text{constraint } i : (b_i/r_i, r_i/n, (b_i - A_i\tilde{\mathbf{x}})/r_i), \quad (10)$$

288 i.e. relaxation value, column degree and fractionality for variables; tight-
 289 ness, scaled row size and relaxation residual for constraints. All are size-
 290 normalized, which is what permits a single parameter set to apply at any
 291 (m, n) .

292 Both node types are embedded to $h = 64$ dimensions by a linear map
 293 and refined by $L = 4$ rounds of bidirectional message passing with graph
 294 attention [18]:

$$\mathbf{h}^{\text{con}} \leftarrow \text{LN}(\text{ELU}(\text{GAT}(\mathbf{h}^{\text{var}} \rightarrow \mathbf{h}^{\text{con}})) + \mathbf{h}^{\text{con}}), \quad (11)$$

$$\mathbf{h}^{\text{var}} \leftarrow \text{LN}(\text{ELU}(\text{GAT}(\mathbf{h}^{\text{con}} \rightarrow \mathbf{h}^{\text{var}})) + \mathbf{h}^{\text{var}}), \quad (12)$$

295 followed by a two-layer head emitting one logit per variable, $\hat{p}_j = \sigma(z_j)$. The
 296 model has 21,761 parameters.

297 The architecture is deliberately small and single-headed. Its only out-
 298 put is the quantity the objective (8) defines; there are no auxiliary value
 299 or feasibility heads, because the search draws soundness from its symbolic
 300 components and asks the network for nothing else.

301 3.7. Inference: guidance, not commitment

302 The network is consulted once per node. Branching selects

$$j^* = \arg \max_{j \in K} |\hat{p}_j - \frac{1}{2}|, \quad \text{first value} = \not\models [\hat{p}_{j^*} \geq \frac{1}{2}], \quad (13)$$

303 that is, the variable the network is most confident about, tried at its pre-
 304 dicted value. The search remains complete and a wrong prediction costs
 305 backtracking only.

306 Variables are *fixed* only by sound deduction. Two mechanisms are avail-
 307 able: constraint propagation, and LP-probing, which fixes x_j when assign-
 308 ing the opposite value makes even the relaxation infeasible. The division

309 of labour follows directly from the measurements in Section 4.5: where a
 310 wrong answer costs time, the cheap unsound predictor is preferable; where
 311 a wrong answer costs correctness, the expensive sound procedure is required
 312 and cannot be replaced.
 313 Algorithm 1 states the resulting search.

Algorithm 1 Guided backtracking search for a BLS

Input: instance $(A, \mathbf{b})$, guidance model f_θ , time budget T
 stack $\leftarrow [(A, \mathbf{b})]$
while stack $\neq \emptyset$ **and** elapsed $< T$:
 $(A', \mathbf{b}') \leftarrow \text{pop}(\text{stack})$
 if $\mathbf{b}' < \mathbf{0}$ **or** $\mathbf{b}' > A'\mathbf{1}$: **continue** // free bound check
 drop all-zero rows of A' ; **if** none remain: **return** any completion
 propagate; **if** contradiction: **continue**
 if propagation determined a set D of variables:
 push reduce($A', \mathbf{b}', D$); **continue** // Proposition 1
 if all variables assigned: **return** solution if verified, else **continue**
 $\hat{\mathbf{p}} \leftarrow f_\theta(A', \mathbf{b}')$ // one forward pass, all variables
 $j^* \leftarrow \arg \max_j |\hat{p}_j - 1/2|$; $v \leftarrow \mathbb{K}[\hat{p}_{j^*} \geq 1/2]$
 push reduce($A', \mathbf{b}', \{j^*=1-v\}$), then reduce($A', \mathbf{b}', \{j^*=v\}$)
return unresolved

314 3.8. Cost-aware fine-tuning of the branching policy

315 The rule of Algorithm 1 branches on the variable the network is surest
 316 about. That maximizes the probability that the first child contains a solution;
 317 it says nothing about how much the search pays when the first child does not,
 318 nor about how much smaller the children are. Both quantities are properties
 319 of the *subtree*, and a search produces them as a by-product: the cost of a
 320 branching decision at node u is the number of nodes expanded below u until
 321 a solution is found there or the subtree is exhausted, and depth-first search
 322 returns it exactly when it backtracks out of u . We use this cost to fine-tune
 323 which variable is branched on.

324 *Decision process..* A state is the reduced instance $(A', \mathbf{b}')$ after propagation;
 325 an action is a free variable j ; the first value tried remains $\mathbb{K}[\hat{p}_j \geq 1/2]$ from
 326 the marginal network, which is not modified. The return of the action at u
 327 is $-\log c(u)$ with $c(u)$ the subtree cost. Because every decision has its own

subtree, credit is assigned per decision rather than per episode, which is what made the sparse-reward formulations we tried earlier unlearnable.

Policy and value.. On the frozen MarginalNet trunk we add two heads. The policy logit is

$$g_j(s) = \log|\hat{p}_j - \frac{1}{2}| + h_\phi(\mathbf{h}_j), \quad (14)$$

where h_ϕ is a two-layer network on the variable embedding $\mathbf{h}_j$ whose last layer is initialized to zero, so that $\pi(j | s) \propto \exp(g_j/\tau)$ starts as a softening of the supervised rule and its argmax is exactly Algorithm 1. The value head $V_\psi(s)$ reads a pooled graph embedding together with $\log |K|$ and $\log m$ and predicts $\log c$.

Update.. Each epoch samples training states, rolls the sampled policy ($\tau = 1$) out to a solution, and records every decision with its exact cost. With advantage $\alpha(u) = V_\psi(s_u) - \log c(u)$, standardized within the epoch, the update is

$$\mathcal{L} = \sum_u \left[-\alpha(u) \log \pi(j_u | s_u) - \lambda \mathcal{H}(\pi(\cdot | s_u)) + (V_\psi(s_u) - \log c(u))^2 \right], \quad (15)$$

with entropy weight $\lambda = 0.01$. Only ϕ and ψ are trained; the trunk and the marginal head are fixed, so the marginal used for the first value and for the analysis of Section 3.4 is unchanged. The critic is pre-trained on decisions recorded under the supervised policy before the actor is updated, and its held-out rank correlation with the true cost is used as a gate: if the baseline cannot rank costs, the advantage is noise and fine-tuning is not attempted. Epoch selection uses the total subtree cost on held-out training-distribution states, never the evaluation instances. At deployment the policy is used greedily: $j^* = \arg \max_j g_j(s)$.

3.9. Cost model

The comparison that motivates the design is per-node cost for a whole-instance verdict, and its structure is asymptotic rather than a constant factor. Let n_c be the number of free variables at a node. The network solves *one* relaxation to build its features and performs *one* forward pass, so its LP count is $O(1)$ in n_c . LP-probing solves one relaxation per candidate variable, $O(n_c)$. The ratio between them therefore grows linearly in instance size, and any single number quoted for it is meaningful only together with the n_c at which it was measured.

359 This also means the comparison is sensitive to how the LP backend is
 360 implemented, in a way that cuts against the learned method. A backend
 361 that rebuilds the model for every call inflates probing’s cost by its full $O(n_c)$
 362 factor while charging the network only once, exaggerating the advantage. A
 363 correctly engineered branch-and-bound keeps one relaxation and re-optimizes
 364 from the parent basis after a bound change, which is an order of magnitude
 365 cheaper per probe. Section 4.5 measures both methods against a warm-
 366 started backend for this reason, and reports the ratio as a function of n_c
 367 rather than as a single figure.

368 4. Experiments

369 4.1. Instances and control of solution multiplicity

370 Instances are generated by the procedure $\text{GENHARD}(m, n, K)$: draw a
 371 random 0/1 matrix A of density $1/2$; plant $K = 20$ candidate solutions
 372 $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(K)}$ with $\mathbf{b}^{(k)} = A\mathbf{x}^{(k)}$; retain the one maximizing *vertex spread*, the
 373 mean pairwise L_2 distance among relaxation vertices obtained from random
 374 objective directions. Maximizing vertex spread enlarges $\tilde{\mathcal{S}}$ directly and is
 375 the mechanism that makes market-split instances hard for relaxation-based
 376 methods (Section 2). Uniqueness of the planted solution is not enforced.

377 *Why $|\mathcal{S}|$ must be controlled across sizes..* Solution multiplicity changes the
 378 *character* of the prediction task, not only its difficulty. When $|\mathcal{S}| = 1$ every
 379 free variable is DET and predicting well is tantamount to solving; when $|\mathcal{S}|$ is
 380 large most variables are OPEN and the task is to estimate a diffuse posterior.
 381 Comparing sizes without controlling $|\mathcal{S}|$ therefore confounds size with task
 382 identity. Since $|\mathcal{S}|$ falls as m rises at fixed n , we select m per size to match
 383 it.

Size	m	m/n	median $ \mathcal{S} $	root ceiling	conditional ceiling
10×25	10	0.40	23	70.4%	87.0%
18×50	18	0.36	19	71.3%	87.2%
21×60	21	0.35	10	—	—

Table 1: Sizes used, with solution multiplicity matched. Ceilings (the Bayes-optimal per-variable accuracies of Section 4.3) agree to within one point between the two training sizes, so size is the variable that changes.

384 *Limits of the control..* Matching $|\mathcal{S}|$ forces m/n to vary ($0.40 \rightarrow 0.35$); at
 385 fixed n the two cannot be held simultaneously and we prioritize $|\mathcal{S}|$ because
 386 it determines task identity. The control is also granular: at $n = 60$ the
 387 median $|\mathcal{S}|$ steps $42 \rightarrow 12 \rightarrow 2$ across $m = 20, 21, 22$, so $m = 21$ is the best
 388 available integer and the realized evaluation set has median $|\mathcal{S}| = 10$ with
 389 27/30 verified.

390 *Where enumeration stops..* At $n = 70$ no enumeration completed within 120s
 391 at the m giving the target multiplicity (0/10). At $n = 100$ the obstruction
 392 is more basic: a constraint solver finds *zero* solutions within 20s at $m \in$
 393 $\{28, 32, 36, 40\}$ despite a planted one. At that size the bottleneck is finding
 394 any solution rather than choosing branches well, so a branching comparison
 395 there would not be informative. We therefore evaluate at $n \leq 60$ and report
 396 this boundary as a limitation.

397 4.2. Implementation and protocol

398 *Software..* Relaxations used for training features and for the original search
 399 loop are solved with SCIP via PySCIPOpt; the warm-started relaxations
 400 of Sections 4.5 and 4.6 use Gurobi 13 (dual simplex, bounds changed in
 401 place), which requires a licence — HiGHS exposes an equivalent incremental
 402 interface and should serve as a free substitute, though we have not verified
 403 the timings with it. Enumeration and the CP-SAT baseline use OR-Tools;
 404 lattice reduction uses fpylll (LLL and BKZ). The network is implemented in
 405 PyTorch with PyTorch Geometric (GATConv). One property of the relaxation
 406 deserves note: it carries no objective, so its solution is an arbitrary vertex
 407 of the relaxation polytope and different LP codes return different vertices.
 408 Both branching arms consume the same vertex within an experiment, so
 409 comparisons are fair, but absolute node counts differ between the SCIP- and
 410 Gurobi-backed runs.

411 *Training..* Adam, learning rate 10^{-3} , 20 epochs over 10,000 in-tree states,
 412 batch size 1 (instances vary in size), gradient-norm clipping at 1.0, loss (8)
 413 with exact marginal targets. Training states are drawn at depths 0–12 uni-
 414 formly, 6 states per source instance, from a pool of 1,760 instances; evaluation
 415 uses 440 instances from a disjoint pool. Training takes approximately 20 min-
 416 utes on one CPU core. *No hyperparameter was tuned*: width, depth, learning
 417 rate, epoch count and the training depth range are the values first tried, and
 418 there is no validation split. Reported figures are therefore not contaminated

by selection on the test set, but neither is any of these settings claimed to be optimal.

Threading.. All timing-sensitive components run single-threaded (`torch.set_num_threads(1)`). On graphs of this size multi-threaded BLAS is counter-productive: we measure 11.3ms per forward pass at 8 threads against 0.9ms at one.

Search evaluation.. Each arm is run on identical instances with a hard per-instance wall-clock budget enforced by subprocess termination, since lattice reduction executes inside a blocking native call that in-process timers cannot interrupt. Instances run at concurrency 6 on a 20-core machine. This low concurrency is deliberate: wall-clock-budgeted comparisons are sensitive to CPU contention, and at higher concurrency we have observed instances solvable in 14–15s standalone exceeding a 20s cap.

Statistics.. Solve-rate differences between arms on the same instances are tested with McNemar’s exact test; per-instance speed differences with a two-sided sign test. We report both rather than a single aggregate because they can have different power. The main comparison at 21×60 is replicated over three independent seeds and we report the spread; smaller-size and transfer results are single-run and marked as such. Seeds are set explicitly: an earlier version derived them from `hash()` of the instance filename, which Python randomizes per process, so those runs had no recorded provenance.

Evaluation sets.. Two 21×60 sets are used: the original 30 instances (generator seeds 700000–700029) and a 100-instance set (seeds 700000–700099) that contains them, built after the 30-instance results had been reported; $|\mathcal{S}|$ is verified for 92/100, median 11. All numbers on the 100-instance set are lattice-off, three seeds. The training pool for the experiments of Sections 4.10 and 4.11 is a separate set of 500 instances at 21×60 (seeds 800000–800499, split 400/100) whose overlap with every evaluation set is zero by hash of $(A, \mathbf{b})$; the 100 held-out instances of that pool supply the states on which fine-tuning epochs are selected.

4.3. Prediction quality against the Bayes ceiling

Because $\mathcal{S}$ is enumerated exactly, the Bayes-optimal per-variable accuracy is computable:

$$\text{ceil} = \frac{1}{|K|} \sum_{j \in K} \max(p_j, 1 - p_j), \quad (16)$$

452 attained by answering the majority value of the surviving solutions. No
 453 deterministic predictor of $\mathbf{x}^*$ from $(A, \mathbf{b})$ can exceed it. Table 2 evaluates
 454 predictors against it on 400 held-out 10×25 instances with $\mathcal{S}$ fully enumer-
 455 ated.

Predictor	all variables	top-3	L_1 to true marginal
Bayes ceiling	70.8%	93.2%	0.0000
Proposed network	66.7%	86.6%	0.1216
LP relaxation rounding	60.9%	65.8%	—

Table 2: Root-level prediction against the exact posterior on the 182 root states of the held-out 10×25 evaluation pool. The network is the same checkpoint used in every search experiment below. “top-3” restricts to the three variables the predictor is most confident about.

456 The network comes within 4.1 points of the ceiling overall and 6.6 on the
 457 top-3 subset, and it recovers most of the gap between relaxation rounding
 458 and the ceiling ($60.9 \rightarrow 66.7$ against a ceiling of 70.8). Two preliminary
 459 experiments on an earlier pool of 10×25 instances, reported for completeness
 460 but not from the checkpoint evaluated here, bear on design choices: a 2.11M-
 461 parameter network with four output heads (selection, assignment, value and
 462 feasibility) scored 60.7% on root states, statistically indistinguishable from
 463 relaxation rounding at 60.6%, so capacity is not the binding constraint and
 464 a single-purpose head is preferable; and training the single-head architecture
 465 on exact marginals rather than the planted solution improved L_1 to the
 466 posterior from 0.1261 to 0.0977 while leaving argmax accuracy unchanged,
 467 which matters because the analysis of Section 3.4 consumes the marginal
 468 rather than its argmax.

469 *What the aggregate conceals..* Table 3 decomposes the same root variables
 470 by the partition of Section 3.4.

471 At the root, 93.7% of variables are OPEN. The aggregate ceiling of 69.9%
 472 is therefore dominated by predictions that cannot affect the search, and on
 473 the 6.3% that can, the information is already complete (100%). Root-level
 474 accuracy is consequently a weak objective for a branching heuristic, and the
 475 operative question is not whether the information exists but what detects it
 476 — taken up next.

Root variables	count	share	ceiling on subset
DET (branching against it empties the subtree)	38	6.3%	100.0%
OPEN (either branch reaches a solution)	562	93.7%	67.2%
all variables	600	100%	69.3%

Table 3: Root-level decomposition, 24 instances of 10×25 with $\mathcal{S}$ enumerated. The aggregate is 94% composed of variables on which no branching decision can cost a backtrack.

4.4. Depth, and where the residual lies

Using Proposition 1 we measure, at each depth along a prefix drawn from an actual solution, the conditional ceiling and the fraction of remaining variables that propagation forces (Table 4).

depth	surviving $ \mathcal{S} $	conditional ceiling	proposed network	LP rounding
0	25.3	70.8%	66.7%	60.9%
4	4.2	82.3%	72.8%	68.5%
7	1.8	91.5%	81.2%	77.8%
8	1.4	94.4%	84.3%	80.9%
12	1.1	98.8%	96.8%	95.5%
all depths	—	87.2%	80.0%	76.5%

Table 4: Conditional information against depth on the held-out in-tree states of the 10×25 evaluation pool ($n = 25$), LP-confidence variable ordering. Same checkpoint as Table 2.

The consequential movement is in composition rather than level: the DET share rises from 6.3% at the root to 87.9% at depth 8 (Table 6). At the root almost every branching decision is inconsequential because both children contain solutions; by depth 8 almost every decision empties a subtree when wrong. The network’s margin over relaxation rounding is 3–6 points throughout and does not vanish at depth, which is the first indication that what it captures is not simply depth-dependent; Section 4.9 makes this precise.

Does in-tree supervision matter, controlling for everything else? A comparison between a root-trained and an in-tree-trained network is only informative if nothing else differs. Table 5 holds the instance pool, the label type (exact marginals), the architecture, the optimizer and the number of training states

492 (1,760 each) fixed, and varies only whether the states are root states or in-
 493 tree states drawn at depths 0–12. Both are evaluated on the same held-out
 494 in-tree states.

depth	conditional ceiling	root-trained	in-tree-trained	LP rounding
0	70.7%	67.8%	66.5%	60.9%
4	82.1%	71.4%	72.4%	68.3%
6	89.7%	74.6%	78.2%	74.0%
8	94.4%	77.4%	84.2%	81.3%
10	97.1%	81.7%	91.4%	89.4%
12	98.8%	85.8%	96.5%	95.3%
all depths	87.2%	75.4%	79.8%	76.4%

Table 5: Controlled comparison of training-state distribution. Same pool, labels, architecture, optimizer and state count; only the depth at which states are sampled differs. Evaluated on held-out in-tree states of the 10×25 pool.

495 In-tree supervision is worth 4.4 points overall. The more telling row is the
 496 root-trained network itself: at 75.4% it is *below* relaxation rounding on in-
 497 tree states, so a network trained on root states does not merely underperform
 498 in the tree, it is worse than the classical heuristic it is meant to replace. The
 499 shift from root to interior states is not a distribution the root-trained network
 500 generalizes across, and in-tree training is what makes the learned guidance
 501 better than the baseline at all.

502 Table 6 decomposes the residual gap at these depths and asks what each
 503 method detects.

depth	% DET	network on DET	propagation	LP integrality	LP-probing	network on OPEN
5	61.9%	89.7%	1.2%	67.9%	51.6%	54.4%
6	74.2%	88.2%	1.6%	70.4%	68.0%	52.5%
7	79.9%	88.8%	3.7%	73.3%	72.8%	51.2%
8	87.9%	89.4%	5.3%	76.9%	79.8%	53.8%

Table 6: Residual decomposition on the same states as Table 4. Columns 4–6 are recall on the ground-truth DET set (LP-probing measured with the warm-started backend of Section 4.5); columns 3 and 7 are the network’s argmax accuracy on each subset.

504 The gap lies entirely on DET variables: at depth 7, $0.80 \times 11.2 \approx 9$ points,
 505 matching the measured 10.3-point shortfall against the conditional ceiling.

506 Performance near 50% on OPEN variables is correct behaviour rather than
507 a defect, since those variables are genuinely undetermined. The striking
508 figure is that unit propagation detects only 1.2–5.3% of logically determined
509 variables: the information is present and derivable but not reachable by the
510 cheapest sound rule. LP-probing recovers 52–80% and LP integrality 68–
511 77%.

512 4.5. Cost per node

513 Both methods are charged everything they must pay to reach a whole-
514 instance verdict: the network one relaxation for its features plus one forward
515 pass, probing one relaxation per free variable. Both use the same warm-
516 started backend, a single model whose variable bounds are changed in place
517 and re-optimized from the parent basis (Gurobi dual simplex), which is how a
518 branch-and-bound implementation reuses its relaxation. Timings are single-
519 threaded.

Instances	n	network	LP-probing	ratio	probing recall
10×25	25	3.22 ms	5.60 ms	$1.7\times$	—
18×50	50	2.33 ms	11.84 ms	$5.1\times$	—
21×60	60	2.57 ms	18.39 ms	$7.2\times$	—
40×100	100	4.02 ms	62.46 ms	$15.5\times$	—
60×150	150	4.45 ms	79.06 ms	$17.8\times$	—
<i>in-tree states of 10×25 instances, by depth:</i>					
depth 5	20	2.44 ms	4.02 ms	$1.6\times$	51.6%
depth 6	19	2.39 ms	3.98 ms	$1.7\times$	68.0%
depth 7	18	2.35 ms	3.43 ms	$1.5\times$	72.8%
depth 8	17	2.34 ms	3.13 ms	$1.3\times$	79.8%

Table 7: Per-node cost for a whole-instance verdict, warm-started LP backend for both methods, single-threaded. The network’s accuracy on DET variables at depths 5–8 is 85.9–92.7% against probing’s recall in the last column.

520 The ratio behaves as the cost model predicts: it is negligible at $n = 25$
521 and reaches $17.8\times$ at $n = 150$, tracking the $O(n_c)$ versus $O(1)$ difference
522 in relaxation count. Two consequences follow, and they point in opposite
523 directions.

524 *The advantage is real but asymptotic, and small where we can search..* At the
525 sizes where our search actually operates ($n = 25\text{--}60$) the ratio is $1.3\text{--}7.2\times$.
526 At the depths where prediction matters most (5–8 of a 10×25 instance,
527 $n_c = 17\text{--}20$) it is $1.3\text{--}1.7\times$. A practitioner choosing between the two at
528 that scale should prefer probing, which is *sound*: where its test fires the
529 conclusion is a proof, whereas the network never offers a guarantee. The
530 learned alternative becomes compelling only at $n \gtrsim 100$, and Section 4.1
531 shows that search itself fails there for every strategy we tried.

532 *The measurement is sensitive to baseline engineering..* An earlier version of
533 this comparison used a backend that rebuilt the LP model on every call
534 and, additionally, excluded the network’s own feature relaxation from its
535 timer. That combination reported $19\text{--}23\times$ at depths 5–8, where the correctly
536 measured value is $1.3\text{--}1.7\times$ — an overstatement of more than an order of
537 magnitude, arising entirely from how the two sides were instrumented rather
538 than from what they compute. We report the corrected figures and note
539 the general hazard: in any cost comparison against a classical procedure, the
540 classical side must be implemented as a practitioner would implement it, and
541 the learned side must be charged for its own preprocessing.

542 4.6. Search performance

543 We compare branching strategies under matched wall-clock budgets with
544 the lattice stage disabled, isolating branching quality. LP selects the variable
545 with the most integral relaxation value and tries that value first; MODEL uses
546 the proposed network. Both run in the same loop: one warm-started relax-
547 ation per node (Section 4.2), identical propagation, identical value ordering,
548 so the arms differ only in which variable they branch on and which value
549 they try first. Each configuration is run three times with different seeds; the
550 depth-first search itself is deterministic given the relaxation vertex, so seeds
551 change timing but not node counts.

552 Three points. First, the learned guidance produces a smaller tree at every
553 size — $2.0\times$, $3.9\times$ and $9.5\times$ fewer nodes than LP — and the ratio grows with
554 size, as the analysis of Section 4.9 predicts: the states where the network
555 separates from relaxation rounding are the large, weakly constrained ones
556 near the top of a big tree, and 10×25 has few of them. Second, a smaller
557 tree buys time only once the tree is large enough to pay for the forward pass.
558 At 10×25 the MODEL arm is *slower* ($0.34\text{--}0.43\times$, faster on at most $1/30$
559 instances), because a 24-node tree costs less than the network evaluations

Size (budget)	strategy	solve rate	median nodes	median time	nodes/s
10×25 (60 s)	LP	$100 \pm 0.0\%$	46	0.03 ± 0.00 s	1,767
	MODEL	$100 \pm 0.0\%$	24	0.07 ± 0.01 s	376
18×50 (300 s)	LP	$100 \pm 0.0\%$	4,224	2.43 ± 0.33 s	2,066
	MODEL	$100 \pm 0.0\%$	1,085	1.82 ± 0.37 s	731
21×60 (600 s)	LP	$100 \pm 0.0\%$	82,366	51.5 ± 2.0 s	1,627
	MODEL	$100 \pm 0.0\%$	8,657	14.6 ± 0.7 s	619

Table 8: Branching strategy comparison on the warm-started loop, lattice stage disabled, 30 instances per cell, mean $\pm$ standard deviation over three seeds of the per-seed median. Node counts are seed-invariant.

that shrink it; its per-node cost is $2.6\text{--}4.7\times$ that of LP at every size, which is the forward pass of Section 4.5. At 18×50 the median per-instance ratio is $1.30\text{--}1.37\times$ in the network’s favour but not significant (faster on 18/30 in every seed, sign test $p = 0.36$). Third, at 21×60 the advantage is clear and replicates: the median per-instance ratio is $2.57\times$, $3.03\times$ and $2.96\times$ across seeds, the network is faster on exactly 23/30 instances in each ($p = 0.005$), and both arms solve every instance within budget.

Two further strategies were measured on an earlier version of this loop that rebuilt the relaxation at every node. RANDOM branching solved 0/30 at 20×50 within 300 s, and Section 4.12 shows it closes none of the 21×60 residual instances on the present loop either. LP-PROBE, which probes every free variable and soundly fixes the forced ones before branching as LP, produced by far the smallest trees ($51\times$ fewer nodes than LP at 20×50) and was the slowest arm, processing 6 nodes/s against 376: node count and wall-clock rank the methods in opposite orders, and a tree-size convention would select the worst of them. We have not repeated LP-PROBE on the warm-started loop; Section 4.5 measures its per-node cost there directly.

Replication on 100 instances.. Thirty instances leave the sign test with little power and the median with wide uncertainty, so the 21×60 comparison was repeated on the 100-instance set (Table 9). The direction and significance hold — MODEL is faster on 70–71/100 instances ($p < 0.001$) with $5.5\times$ fewer nodes — but the margins are smaller than on the original thirty: $2.0\text{--}2.3\times$ in time against $2.6\text{--}3.0\times$, and $5.5\times$ in nodes against $9.5\times$. The original thirty are the first thirty seeds of the larger set, and on them the supervised

584 policy happens to do unusually well; we treat the 100-instance figures as
585 the estimate and the 30-instance ones as the favourable end of its sampling
586 variation.

strategy	solve rate	median nodes	median time	nodes/s	vs. LP (per instance)
LP	100 ± 0.0%	79,609	41.2 ± 1.9 s	1,966	—
MODEL	100 ± 0.0%	13,386	16.5 ± 0.6 s	796	2.0–2.3×, 70–71/100, $p < 0.001$
MODEL+FT (Section 4.11)	100 ± 0.0%	5,495	7.6 ± 0.3 s	738	4.8–5.9×, 81–84/100, $p < 0.001$

Table 9: 21×60 , 100 instances, lattice off, 600 s budget, three seeds (mean ± s.d. of per-seed medians; node counts are seed-invariant). The last row is the cost-aware fine-tuned policy of Section 4.11.

587 We also record what the rebuilt-relaxation loop reported for the 21×60
588 comparison, because an earlier draft of this paper rested on it: 30/30 against
589 25/30 solved, and a 3.29× median speedup. Warm-starting the relaxation
590 removed the solve-rate difference entirely — the five instances LP failed within
591 600 s it now closes in 65–270 s — and reduced the speedup. The solve-rate
592 advantage was an artifact of a slow backend that penalized the arm expanding
593 more nodes, and we do not claim it.

594 4.7. Comparison against complete solvers

595 Table 8 compares branching strategies inside our own search loop, which
596 measures branching quality but says nothing about how the system stands
597 against an off-the-shelf solver. That comparison is reported here, on the same
598 instances, single-threaded, under the same budget.

Size	CP-SAT	SCIP	proposed (MODEL)	CP-SAT advantage
10×25	0.002 s	0.009 s	0.07 s	35×
18×50	0.046 s	0.385 s	1.82 s	40×
21×60	1.159 s	3.687 s	14.6 s	13×
21×60 , 100 inst.	0.792 s	2.590 s	16.5 s / 7.6 s (fine-tuned)	21× / 10×

Table 10: Median solve time; all methods solve every instance at every size. The proposed column is the MODEL arm of Table 8 (mean of per-seed medians); the last row adds the 100-instance set and the fine-tuned policy of Section 4.11. The proposed system does not beat a complete solver at any size tested, and is outperformed by SCIP, whose LP code it uses for training features.

599 The gap is large and we state it plainly: *the contribution of this paper is*
600 *not a solver that competes with CP-SAT*. It is worth asking where the gap

comes from, because the answer bears on what the branching heuristic is actually doing. Table 11 decomposes it at 21×60 into tree size and per-node throughput, using CP-SAT’s own branch counter.

	CP-SAT	proposed	ratio
nodes (median)	18,827 branches	8,657	ours $2.2\times$ fewer
throughput	19,782 nodes/s	619 nodes/s	CP-SAT $32\times$ faster
total time (median)	1.065 s	14.6 s	$14\times$

Table 11: Decomposition of the 21×60 gap. CP-SAT additionally learns a median of 5,514 conflict clauses per instance.

The learned guidance produces a *smaller* search tree than CP-SAT — $2.2\times$ fewer nodes despite CP-SAT’s clause learning — and loses entirely on per-node cost. Where that cost goes can be measured. Before the loop was warm-started it processed 254 nodes/s, spending 4.34 ms per node rebuilding and solving the relaxation; warm-starting the same LP brings that to 0.08 ms and raises throughput to 619 nodes/s, which is the loop reported throughout Section 4.6. Of the 1.6 ms that remain per node, the network’s forward pass is 1.88 ms at the root and less at deeper nodes where the reduced graph is smaller, and propagation and bookkeeping are under 0.1 ms: the forward pass now accounts for over 90% of per-node cost, and we have not reduced it. We therefore do not claim that the gap is an implementation artifact. What the node counts establish is narrower: the deficit is not in branching quality. Closing it would require at minimum a compiled inference path and, as the next paragraph notes, algorithmic components our search lacks.

Two honest caveats attach to this reading. Tree sizes are not perfectly commensurable — CP-SAT’s “branches” include restarts and conflict-driven jumps, and each of its nodes performs more propagation work than ours — so the $2.2\times$ should be read as an order-of-magnitude statement. And a competitive implementation would need more than a faster forward pass: clause learning and restarts are absent from our search, and the last two are what let CP-SAT scale past the sizes where we stop.

4.8. Size transfer

Because the features are size-normalized and the parameters do not depend on (m, n) , one network applies at any size. We test whether that transfers in practice, with $|\mathcal{S}|$ matched as in Table 1.

Training size	Evaluation size	solve rate	median time	median nodes
10×25	10×25	100%	0.07 ± 0.01 s	24
10×25	18×50	100%	1.82 ± 0.37 s	1,085
18×50	18×50	100%	1.48 ± 0.15 s	942
10×25	21×60	100%	14.6 ± 0.7 s	8,657
18×50	21×60	100%	23.7 ± 0.7 s	14,742

Table 12: Transfer on the warm-started loop, lattice stage disabled, 30 instances per row, three seeds. Rows 2–3 and 4–5 each hold the evaluation size fixed and vary only the training size.

At the 21×60 evaluation size the network trained only at 10×25 — a $2.4 \times$ smaller problem — is not worse than the one trained at 18×50 ; if anything it is better, expanding fewer nodes on $18/30$ instances and finishing faster on $18\text{--}19/30$ in each seed, with a median per-instance time ratio of $1.24\text{--}1.31 \times$ in its favour. Neither difference is significant (sign test $p \geq 0.20$), and both solve $30/30$. At 18×50 the picture is the same: the 10×25 network is faster on $16\text{--}17/30$ instances ($p \geq 0.59$) although its median time is slightly higher, the two summaries disagreeing because the paired differences are small against instance-to-instance variation. A previous version of this comparison, on the rebuilt-relaxation loop, gave $15/30$ and a $0.92 \times$ ratio; the direction is not stable across loops and we read the result as no detectable transfer loss rather than as a transfer gain. Root-level prediction agrees: the 10×25 network scores 67.8%, 69.5% and 69.4% at 10×25 , 20×50 and 40×100 , holding its margin over relaxation rounding (+11.0, +7.2, +4.7 points) at every size.

This matters practically because label generation is the binding constraint. Exact marginals require enumeration, which fails beyond $n \approx 60$; transfer means a network trained where labels are obtainable can be deployed where they are not.

4.9. Where the transferred network helps

Section 4.4 located the useful prediction regime at depths 6–8 of a 10×25 instance. Whether that localization survives transfer is a separate question, and it does not: at 21×60 the search queries the network on states with 5–60 free variables, only 19.7% of which fall inside the training range of $13\text{--}25$, and 80.2% of which are larger than any training state. Table 13 evaluates

the network and relaxation rounding against exact conditional marginals on 21×60 in-tree states, grouped by the number of free variables, on the subset of states that could be enumerated within 15 s.

free vars	in training range	$m/ K $	network on DET	LP on DET	gain	corr(network, LP)
20	yes	1.05	100.0%	100.0%	+0.0	1.000
25	yes	0.84	100.0%	100.0%	+0.0	0.998
30	no	0.70	95.7%	95.0%	+0.7	0.954
40	no	0.53	75.5%	75.2%	+0.4	0.747
50	no	0.42	70.2%	65.0%	+5.3	0.664
60	no	0.35	93.3%	88.9%	+4.4	0.658

Table 13: 21×60 in-tree states by number of free variables, 9–14 enumerable states per row. Gain is the network’s margin over relaxation rounding on DET variables; the last column is the Pearson correlation between the two predictors’ outputs.

The pattern is the reverse of what a depth-based account predicts. Inside the training range the network is indistinguishable from relaxation rounding — both are perfect and their outputs correlate at 1.000 — because at $m/|K| \geq 0.8$ the reduced instance is over-constrained and the relaxation is already integral on every DET variable. The network’s entire margin arises on states it never saw during training, at 50–60 free variables, where the relaxation is weakest ($m/|K| \approx 0.35$ – 0.42) and its outputs have decorrelated from the relaxation’s. In terms of Section 3.4, this is the region where DET variables are present but the relaxation does not find them.

Two conclusions follow. The first corrects Section 4.4: depth was a proxy for constraint density that happened to hold on 10×25 , where descending from the root moves $m/|K|$ only from 0.40 to 0.59, and does not hold on 21×60 , where the same descent moves it from 0.35 past 1.0. The useful regime is defined by $m/|K|$, and at deployment size it sits near the *top* of the tree. The second is that the transfer result of Section 4.8 is stronger than a matched-multiplicity comparison alone shows: the network extrapolates to states outside its training range in both size and constraint density, and that extrapolation is where its contribution lives. It is also a limitation, since the training range was not chosen with this in mind (Section 4.14).

4.10. Training on deployment-sized states: a negative result

Section 4.9 located the network’s margin on states with 50–60 free variables, outside the 13–25 range of its training states, and Section 4.14 of an

earlier draft listed aligning the two as an obvious untested improvement. We tested it, holding everything but the training-state distribution fixed.

States.. From a separate 21×60 pool (Section 4.2) we built in-tree states with 30, 35, ..., 60 free variables by fixing a prefix of the LP-confidence ordering to the planted solution, and labelled them with exact conditional marginals under the enumeration rule of Section 3.5. Every state with at most 55 free variables enumerated within 90 s; at 60 (the root) 291/400 did, and the failures are the instances with large $|\mathcal{S}|$, so root labels are biased toward small solution sets. A property of these states matters for what follows: with the prefix fixed to a solution, the surviving solution set is almost always a singleton once 50 or fewer variables remain (median $|\mathcal{S}(F, \mathbf{v})| = 1$, DET share 97%), so their labels are nearly hard. Three networks were trained with the architecture, optimizer, epochs and loss of Section 4.2: M0, the network used everywhere above (10×25 in-tree states); M1, the 2,691 deployment-range states only; M2, half of each. A pre-diagnostic on 10×25 alone, training on depth ranges 0–6, 6–12 and 0–12 with equal state counts, gave in-tree accuracies of 79.3%, 79.9% and 79.9%: within one size, the depth range does not matter.

Prediction.. Re-running the analysis of Table 13 with the three networks scored on identical states, the gain over relaxation rounding on DET variables at 50 and 60 free variables is +5.4 and +9.7 points for M0, +6.1 and +10.7 for M1, +5.9 and +10.6 for M2. Training in the deployment regime buys about one point there.

Search.. It costs far more elsewhere. On the 100-instance set M1 expands a median of 21,337 nodes against 13,386 for M0 ($1.6\times$ more; fewer on only 38/100, $p = 0.021$) and is $1.7\text{--}1.8\times$ slower; on the original thirty it expands 20,746 against 8,657. M2 sits between (18,558, $p = 0.089$). A switch that routes nodes with ≥ 50 free variables to M1 and the rest to M0 recovers nothing (17,252 nodes, 43/99, $p = 0.23$). The mechanism is visible in the prediction diagnostics: on small, dense interior states (10×25 at depth 6–8, the kind the search spends two thirds of its nodes on) the probability that M1’s most confident variable is DET and predicted correctly falls from 94.4% to 77.6%, and the rank correlation between its confidence and the true margin from 0.36 to 0.12, while on the deployment-range states it improves ($77.2 \rightarrow 86.0\%$). Nodes are where the interior is, and a network that is slightly better near the top of the tree and much worse inside it makes the

715 search slower. We conclude that the training-state size distribution is not the
 716 lever, and that the residual identified in Section 4.9 is not a distribution-shift
 717 artifact.

718 4.11. Cost-aware fine-tuning

719 If neither accuracy on DET variables nor the training distribution explains
 720 what the supervised guidance leaves on the table, the natural suspect is the
 721 objective: the rule branches on the variable it is surest about, and being sure
 722 is not the same as being cheap. Three measurements test this, and the third
 723 turns it into a method.

724 *R0: does the confident choice cost the least?* On 100 held-out states with
 725 45–50 free variables we forced the root decision to each of the five variables
 726 the network is most confident about, ran the supervised search below it,
 727 and recorded the exact subtree cost of each. The confidence rule’s choice
 728 is the cheapest of the five on 22% of states; the median ratio of its cost to
 729 the cheapest candidate is 1.29 and the mean 2.97; on 24% of states some
 730 candidate is at least twice as cheap; and choosing the cheapest candidate at
 731 the root alone would remove 55% of all nodes summed over these states. The
 732 LP rule’s root choice costs the same as the network’s on only 10% of states,
 733 so the two rules are not merely tied. There is a large residual, and it is about
 734 *which* variable to branch on.

735 *R1: can the cost be predicted?* The critic is the baseline of the policy update,
 736 so it must rank decisions by cost on states it has not seen. Recording every
 737 decision of the supervised search on the training pool gave 37,275 decisions
 738 (9,375 on held-out states) with their exact subtree costs, of which two thirds
 739 occur at 25–34 free variables. A value head on the frozen trunk reaches a
 740 held-out Spearman correlation of 0.80 with $\log c$; a baseline that knows only
 741 the state’s size (a linear fit on $\log |K|$) reaches 0.70, and within size buckets
 742 the critic’s advantage is larger (0.68 against 0.38 at 35–44 free variables,
 743 0.57 against 0.37 at 45–55). Unfreezing the trunk adds 0.02. The gate of
 744 Section 3.8 is passed with the frozen critic.

745 *R2: fine-tuning.* Six epochs of 400 rollouts each ($\tau = 1$, 90s budget, about
 746 21,000 decisions per epoch) from the 45–60 free-variable training states. On
 747 the 150 held-out monitor states the total subtree cost under the greedy policy,
 748 relative to the supervised rule, was 0.79, 0.49, 0.56, 0.46, 0.65 and 0.56 after

epochs 1–6; epoch 4 was selected. Policy entropy fell from 1.73 to 0.10 over the same epochs, so the policy is near-deterministic by the selected epoch.

Evaluation. Table 9 reports the fine-tuned policy (MODEL+FT) beside the supervised one on the 100-instance set: 5,495 nodes against 13,386 ($2.3\times$ fewer, fewer on 70/100, $p < 0.001$ in every seed) and a median of 7.6 s against 16.5 s ($2.0\text{--}2.2\times$, faster on 64–67/100, $p \leq 0.007$), at an unchanged per-node cost (738 against 796 nodes/s). Against LP-guided branching this is $4.8\text{--}5.9\times$ in time and $14\times$ in nodes. On the original thirty instances the effect is small (8,873 against 8,657 nodes; $1.2\text{--}1.5\times$ in time, 18–21/30, p between 0.04 and 0.36): the seventy added instances account for the gain ($2.5\times$ fewer nodes, 53/70), which is consistent with the thirty being a favourable sample for the supervised rule. In the deployed cascade (Table 14) the fine-tuned policy closes the lattice residual in 7.3–11.7 s against 12.2–17.7 s for the supervised one.

What changed. On 300 held-out states the fine-tuned policy branches on a different root variable than the confidence rule at 95% of them. Its pick is typically the 15th most confident variable ($|\hat{p} - \frac{1}{2}|$ of 0.26 against 0.43 for the rule’s pick), and it is DET with a correct first value less often (57.5% against 82.5%). The policy has moved *away* from the variables it is surest about, and the search got cheaper. The improvement is not attributable to the root decision alone: forcing only the root to the fine-tuned pick and running the supervised policy below it gives a median cost ratio of 0.93 (cheaper on 22/46 disagreement states), so the gain is a property of the sequence of decisions rather than of any single one. We can therefore say what the fine-tuning optimizes — subtree cost, with the probability of being right held as a soft prior — but not yet name the feature it found; a cheap proxy, the amount of propagation the first child triggers, does not differ between the two choices.

4.12. Deployment in the full pipeline

Table 8 isolates branching by disabling lattice reduction. In deployment the lattice stage runs first and the search sees only its residual. Because that stage is independent of the branching heuristic, every arm receives the identical residual set.

Lattice coverage falls with size — 30/30, 24–28/30, 10–16/30 — and that decline is what creates room for branching quality to matter. At 10×25 and 18×50 the residual is empty or closed within seconds by either sound

Size	strategy	closed by lattice	residual solved by search	median time on residual	end-to-end
10×25	both	30/30	—	—	100%
18×50	LP / MODEL	24–28/30	11/11 / 11/11	0.5–5.0 s / 0.6–2.5 s	100% / 100%
21×60	RANDOM	10–16/30	0/50	> 600 s	$44.4 \pm 8.3\%$
21×60	LP	10–16/30	50/50	52.9–65.1 s	$100 \pm 0.0\%$
21×60	MODEL	10–16/30	50/50	14.2–20.2 s	$100 \pm 0.0\%$
21×60	MODEL+FT	10–16/30	50/50	7.3–11.7 s	$100 \pm 0.0\%$

Table 14: Cascade results on the warm-started loop, three seeds at every size, 600s budget at 21×60 . Lattice coverage varies across seeds (16, 14, 10 of 30 at 21×60) because the lattice stage draws random column permutations, so the residual handed to the search varies with it; all arms within a seed receive the identical residual.

strategy. At 21×60 between ten and twenty instances reach the search depending on the seed, and the residual is genuinely hard: random branching closes none of the 50 residual instances pooled over seeds, so the end-to-end rate of that arm is the lattice coverage alone. Both relaxation-based and learned branching close all 50, and the difference between them is time. On the residual the learned guidance is faster on 11/14, 14/16 and 15/20 instances in the three seeds (median per-instance ratio $3.48\times$, $2.99\times$, $2.38\times$; sign test $p = 0.057$, 0.004 , 0.041), 40/50 pooled with a median ratio of $3.00\times$ ($p = 2.4 \times 10^{-5}$), expanding 6.9–9.5 $\times$ fewer nodes. The residual is where the branching heuristic operates in deployment, and the speedup there matches the isolated measurement of Table 8.

4.13. Does the graph structure contribute?

MarginalNet consumes six scalars per node and propagates them over the bipartite graph. With a feature set that small, a per-variable model ignoring the graph is a serious baseline: if it matches, the message passing is decoration. Table 15 compares models trained on identical targets, loss and splits, differing only in whether structure is used. The aggregated variant appends mean and max of the constraint-side features over each variable’s incident rows, which is the information one message-passing round would deliver.

Message passing contributes: +4.7 points over the best structure-free model and a 36% reduction in L_1 , closing roughly 60% of the residual gap to the ceiling that the MLP leaves open. The improvement in L_1 is the more relevant of the two, since the analysis of Section 3.4 consumes the marginal rather than its argmax.

Model	accuracy	L_1 to marginal	parameters
Linear, variable features only	78.1%	0.2530	4
MLP, variable features only	77.8%	0.2522	8,641
MLP + constraint aggregates	79.4%	0.2361	9,025
MarginalNet (bipartite graph)	84.1%	0.1514	21,761
Bayes ceiling	87.3%	0.0000	—

Table 15: Ablation on in-tree states of 10×25 instances, 2,000 training and 800 test states.

809 The same table bounds the claim, though. A four-parameter linear model
810 on the relaxation value alone reaches 78.1%, within 9.2 points of the ceiling;
811 most of the available signal is in that single feature, and the graph adds a
812 useful but not dominant increment on top of it.

813 4.14. Limitations

814 *Label generation bounds the method, not only its evaluation..* Exact marginals
815 require enumerating $\mathcal{S}$, which does not complete beyond $n \approx 60$ on this fam-
816 ily. Transfer permits deployment at larger sizes but not training there; ap-
817 proximating p_j by solution sampling is the natural extension and is untested.

818 *The useful size window is bounded on both sides..* Below it, lattice reduction
819 closes everything and branching quality is irrelevant. Above it ($n = 100$),
820 search fails outright for every strategy, since a constraint solver finds no
821 solution within 20 s. Our demonstration lives at $n = 60$ and we do not show
822 the window extends further.

823 *The advantage is in time and tree size, not in solve rate..* On a warm-started
824 loop both sound strategies solve every instance we evaluate within budget;
825 what the supervised guidance buys at 21×60 is a $2.0\text{--}2.3\times$ median speedup
826 and $5.5\times$ fewer nodes on 100 instances ($2.6\text{--}3.0\times$ and $9.5\times$ on the original
827 thirty, a favourable sample), and the fine-tuned policy a further $2.0\text{--}2.2\times$
828 and $2.3\times$, significant per instance in every seed. An earlier draft reported
829 a 5–0 solve-rate advantage from a loop that rebuilt the relaxation at every
830 node; that difference vanished with the faster backend and we do not claim it.
831 Whether a solve-rate advantage exists at larger budgets or sizes is untested,
832 and with 30 evaluation instances per size the power to detect one would be
833 low in any case.

834 *Multiplicity control is approximate..* Median $|\mathcal{S}|$ is 23, 19 and 10 across the
 835 three sizes; at $n = 60$ multiplicity steps $42 \rightarrow 12 \rightarrow 2$ across consecutive inte-
 836 ger m , so finer matching is unavailable. Three of thirty evaluation instances
 837 have unverified $|\mathcal{S}|$ and are marked as such. Matching $|\mathcal{S}|$ also moves m/n
 838 from 0.40 to 0.35.

839 *Cost advantage depends on scale and on baseline engineering..* The amortiza-
 840 tion argument holds asymptotically but is worth little at the sizes where our
 841 search operates: $1.3\text{--}1.7\times$ at the depths where prediction matters, against a
 842 warm-started probing baseline that is additionally sound. It reaches $15\text{--}18\times$
 843 only at $n \geq 100$, where search fails for every strategy we tested. We therefore
 844 cannot demonstrate the regime in which the argument is strongest.

845 *The search implementation is not competitive..* At 21×60 CP-SAT is $9\text{--}10\times$
 846 faster than the fine-tuned policy and $20\times$ faster than the supervised one on
 847 the 100-instance set ($13\times$ on the original thirty). Our branching produces
 848 fewer nodes, so the deficit is per-node throughput ($32\times$), now dominated by
 849 the network’s forward pass; we have not shown that a competitive implemen-
 850 tation is achievable: a compiled inference path, clause learning and restarts
 851 are all absent, and the latter two are load-bearing for CP-SAT at larger sizes.

852 *Training and deployment distributions barely overlap, and aligning them did*
 853 *not help..* The training range $\{0, \dots, 12\}$ produces states with 13–25 free
 854 variables, whereas the 21×60 search queries states with 5–60 and spends
 855 80% of its queries above 25. Section 4.10 shows that training on the deploy-
 856 ment regime makes the search slower; we do not know whether a different
 857 construction of large training states (prefixes from the search’s own trajec-
 858 tory rather than from a solution, or states with several surviving solutions)
 859 would behave differently.

860 *The fine-tuning result rests on one training pool and few training seeds..* The
 861 policy was fine-tuned on states from one 21×60 pool and evaluated at that
 862 size only

863 *Hyperparameters are untuned and there is no validation split for the super-*
 864 *vised network..* See Section 4.2. The configuration is defensible as a first
 865 attempt and is not contaminated by the test set; the depth range, at least,
 866 was subsequently shown not to matter within a size (Section 4.10). The fine-
 867 tuning stage does select an epoch on held-out states, which are drawn from
 868 the training pool and disjoint from every evaluation instance.

869 *Single instance family.* All results concern one generator from one appli-
 870 cation. Whether the amortization argument transfers to other constraint
 871 families is untested.

872 5. Conclusion

873 We presented a learned branching heuristic for binary linear systems to-
 874 gether with an exact procedure for supervising it. Two elements carry the
 875 method. Conditioning a BLS on a partial assignment is identical to reducing
 876 it, so a single network trained on ordinary instances predicts at every node of
 877 the search tree without architectural change. And on instances small enough
 878 to enumerate, the exact conditional posterior both supplies the training tar-
 879 get and identifies which predictions can affect the search at all: only variables
 880 that every surviving solution agrees on can cost a backtrack, and their share
 881 rises from 6.3% at the root to 87.9% at depth 8.

882 Measured inside our search loop the heuristic does what it is designed
 883 to do. It is $2.0\text{--}2.3\times$ faster in median than relaxation-based branching at
 884 21×60 on 100 instances in each of three seeds, expands $5.5\times$ fewer nodes,
 885 and closes the residual of the deployed pipeline $3.0\times$ faster, where random
 886 branching closes nothing; with solution multiplicity controlled, a network
 887 trained at $2.4\times$ smaller scale performs indistinguishably from one trained at
 888 the evaluation size, which matters because exact supervision is available only
 889 at small sizes.

890 We are equally explicit about what the results do not establish. The sys-
 891 tem does not compete with a complete solver: CP-SAT is $9\text{--}40\times$ faster at
 892 every size we tested. Decomposing that gap is informative rather than excul-
 893 patory — our branching yields $2.2\times$ fewer nodes and loses on throughput by
 894 $32\times$, which places the deficit in a Python research implementation whose per-
 895 node cost is now the network’s forward pass, and which lacks clause learning
 896 and restarts — but we have not shown that closing it is achievable. Nor does
 897 the cost argument carry the weight we initially assigned it. Against a prop-
 898 erly warm-started probing baseline the advantage is $1.3\text{--}1.7\times$ at the depths
 899 where prediction matters on our instances, and probing is sound where the
 900 network is not; the ratio reaches $15\text{--}18\times$ only at $n \geq 100$, a regime in which
 901 search fails for every strategy we tried. An earlier version of this compar-
 902 ison, using a backend that rebuilt the LP for each probe and excluded the
 903 network’s own feature solve from its timer, reported $19\text{--}23\times$ at a scale where
 904 the true figure is near unity, and we record the correction because the error

905 came entirely from instrumentation rather than from what the two methods
906 compute.

907 Two further results sharpen what the supervised target is. Moving the
908 training states to the deployment regime, where the network’s prediction
909 margin lives, does not help and in fact doubles the search tree, because the
910 search spends its nodes in the small interior states the moved network no
911 longer handles well; the training-state distribution is not the lever. Optimiz-
912 ing the cost the search pays is. With the exact subtree cost of every decision
913 read off the search, an actor-critic fine-tuning of the branching choice —
914 initialized at the supervised rule, with the marginal network frozen — ex-
915 pands $2.3\times$ fewer nodes and runs $2.0\text{--}2.2\times$ faster than the supervised policy
916 on 100 instances, and the policy it finds branches on variables the network
917 is *less* sure about. The probability that a branch is right and the cost of
918 the branch being wrong are different objectives, and the second is the one a
919 search minimizes.

920 What remains, and what we claim, is methodological. For planted-
921 solution instance families the exact conditional posterior is computable and is
922 a well-posed supervision signal for branching, where the conventional target
923 — a single planted solution — is not. States drawn from inside the search
924 tree are a better training distribution than root states, and the reduction
925 identity makes generating them a matter of substitution rather than solver
926 instrumentation. The resulting heuristic transfers across instance size when
927 solution multiplicity is held comparable, and the exact cost of its decisions,
928 available from the search, is a second well-posed training signal that im-
929 proves it. Whether this becomes practically useful depends on two things
930 we have not done: extending exact supervision past $n \approx 60$, most plausibly
931 by approximating the marginals without exhaustive enumeration, and imple-
932 menting the search well enough that a branching improvement is visible at
933 the system level.

934 Reproducibility

935 Code is organized by component under `proposed_src/`, with shared mod-
936 ules in `util/` and a `ROLES.txt` in each directory describing every file. Scripts
937 are run from their own directory. Every search number in Tables 8, 12
938 and 14 comes from one chain, `v23_ablation/run_revision2.sh`, which runs
939 the warm-started loop (`v23_warm_search.py`) over all conditions and seeds;
940 `v23_aggregate.py` produces the tables from its output.

```

941 # instances with matched |S| and exact marginals
942 cd proposed_src/v15_bayes
943 python3 v15_gen_by_m.py --m 10 --n 25 --count 2200 --K 20 \
944     --out ../../runs/v22/sols_10x25.json
945
946 # attach solution sets, then split 80/20 (seed 0)
947 cd ../v22_transfer
948 python3 v22_add_solutions.py \
949     --inp ../../runs/v22/sols_10x25.json \
950     --out ../../runs/v22/full_10x25.json
951 python3 v22_split.py # -> full_10x25_{train,test}
952
953 # in-tree training states (depths 0-12, 6 per instance)
954 cd ../v17_conditional
955 python3 v17_make_conditional_data.py \
956     --sols ../../runs/v22/full_10x25_train.json \
957     --max_depth 12 --per_instance 6 \
958     --out ../../runs/v22/cond_10x25_train.json
959
960 # train (about 20 minutes, single CPU core)
961 python3 v17_train_conditional.py \
962     --train ../../runs/v22/cond_10x25_train.json \
963     --test  ../../runs/v22/cond_10x25_test.json \
964     --n_train 10000 --n_test 2000 --epochs 20 \
965     --target soft --out ../../runs/model_n25
966
967 # every search table: warm-started loop, 3 seeds,
968 # conditions A/B/R/C/D/Crand x {lattice off, on}
969 cd ../v23_ablation
970 bash run_revision2.sh # -> runs/v23/, chain.log
971 python3 v23_aggregate.py # -> runs/v23/aggregate.json
972
973 # v24: range training (negative), 100-instance
974 # set, R0/R1/R2 fine-tuning and evaluation
975 cd ../v24_deploy_range
976 bash run_v24.sh; bash run_v24_switch.sh
977 bash run_v24_r1.sh; bash run_v24_r2.sh
978 python3 v24_aggregate.py \
979     --tags M0 M1 M2 SW PV
980
981 All generators are seeded. The released checkpoints are 96 KB each. In-

```

981 stance data and weights are regenerated from the above rather than dis-
982 tributed, and the full evaluation cycle is scripted in `run_revision2.sh` under
983 `v23_ablation/`.

984 References

- 985 [1] Aardal, K., Hurkens, C. A. J., and Lenstra, A. K. (2000). Solving a
986 system of linear Diophantine equations with lower and upper bounds on
987 the variables. *Mathematics of Operations Research*, 25(3):427–442.
- 988 [2] Aardal, K., Bixby, R. E., Hurkens, C. A. J., Lenstra, A. K., and
989 Smeltink, J. W. (2000). Market split and basis reduction: Towards a
990 solution of the Cornuéjols–Dawande instances. *INFORMS Journal on*
991 *Computing*, 12(3):192–202.
- 992 [3] Amos, B. (2023). Tutorial on amortized optimization. *Foundations and*
993 *Trends in Machine Learning*, 16(5):592–732.
- 994 [4] Bengio, Y., Lodi, A., and Prouvost, A. (2021). Machine learning for
995 combinatorial optimization: A methodological tour d’horizon. *European*
996 *Journal of Operational Research*, 290(2):405–421.
- 997 [5] Chen, Z., Liu, J., Wang, X., Lu, J., and Yin, W. (2023). On represent-
998 ing mixed-integer linear programs by graph neural networks. In *Inter-*
999 *national Conference on Learning Representations*.
- 1000 [6] Cornuéjols, G. and Dawande, M. (1999). A class of hard small 0-1 pro-
1001 grams. In *Integer Programming and Combinatorial Optimization*, pages
1002 284–293.
- 1003 [7] Gasse, M., Chételat, D., Ferroni, N., Charlin, L., and Lodi, A. (2019).
1004 Exact combinatorial optimization with graph convolutional neural net-
1005 works. In *Advances in Neural Information Processing Systems*.
- 1006 [8] Gershman, S. J. and Goodman, N. D. (2014). Amortized inference in
1007 probabilistic reasoning. In *Proceedings of the 36th Annual Meeting of*
1008 *the Cognitive Science Society*, pages 517–522.
- 1009 [9] Khalil, E. B., Le Bodic, P., Song, L., Nemhauser, G., and Dilkina, B.
1010 (2016). Learning to branch in mixed integer programming. In *AAAI*
1011 *Conference on Artificial Intelligence*, pages 724–731.

- 1012 [10] Lenstra, A. K., Lenstra, H. W., and Lovász, L. (1982). Factoring poly-
1013 nomials with rational coefficients. *Mathematische Annalen*, 261(4):515–
1014 534.
- 1015 [11] Li, Z., Guo, J., and Si, X. (2023). G4SATBench: Benchmarking and
1016 advancing SAT solving with graph neural networks. *Transactions on*
1017 *Machine Learning Research*.
- 1018 [12] Marques-Silva, J. P. and Sakallah, K. A. (1999). GRASP: A search algo-
1019 rithm for propositional satisfiability. *IEEE Transactions on Computers*,
1020 48(5):506–521.
- 1021 [13] Millidge, B. (2022). Deconfusing direct vs amortized opti-
1022 mization. Unrefereed technical note, [https://www.beren.io/](https://www.beren.io/2022-09-25-Deconfusing-direct-vs-amortized-optimization/)
1023 [2022-09-25-Deconfusing-direct-vs-amortized-optimization/](https://www.beren.io/2022-09-25-Deconfusing-direct-vs-amortized-optimization/).
- 1024 [14] Moskewicz, M. W., Madigan, C. F., Zhao, Y., Zhang, L., and Malik, S.
1025 (2001). Chaff: Engineering an efficient SAT solver. In *Design Automa-*
1026 *tion Conference*, pages 530–535.
- 1027 [15] Ohrimenko, O., Stuckey, P. J., and Codish, M. (2009). Propagation via
1028 lazy clause generation. *Constraints*, 14(3):357–391.
- 1029 [16] Schnorr, C. P. and Euchner, M. (1994). Lattice basis reduction: Im-
1030 proved practical algorithms and solving subset sum problems. *Mathe-*
1031 *matical Programming*, 66(1–3):181–199.
- 1032 [17] Selsam, D., Lamm, M., Bünz, B., Liang, P., de Moura, L., and Dill,
1033 D. L. (2019). Learning a SAT solver from single-bit supervision. In *In-*
1034 *ternational Conference on Learning Representations*.
- 1035 [18] Veličković, P., Cucurull, G., Casanova, A., Romero, A., Liò, P., and Ben-
1036 gio, Y. (2018). Graph attention networks. In *International Conference*
1037 *on Learning Representations*.
- 1038 [19] Williams, R. J. (1992). Simple statistical gradient-following algorithms
1039 for connectionist reinforcement learning. *Machine Learning*, 8(3–4):229–
1040 256.
- 1041 [20] Etheve, M., Alès, Z., Bissuel, C., Juan, O., and Kedad-Sidhoum, S.
1042 (2020). Reinforcement learning for variable selection in a branch and

- 1043 bound algorithm. In *Integration of Constraint Programming, Artificial*
1044 *Intelligence, and Operations Research (CPAIOR)*, pages 176–185.
- 1045 [21] Scavuzzo, L., Chen, F. Y., Chételat, D., Gasse, M., Lodi, A., Yorke-
1046 Smith, N., and Aardal, K. (2022). Learning to branch with tree MDPs.
1047 In *Advances in Neural Information Processing Systems*, volume 35.
- 1048 [22] Haralick, R. M. and Elliott, G. L. (1980). Increasing tree search efficiency
1049 for constraint satisfaction problems. *Artificial Intelligence*, 14(3):263–
1050 313.
- 1051 [23] Xu, K., Hu, W., Leskovec, J., and Jegelka, S. (2019). How powerful
1052 are graph neural networks? In *International Conference on Learning*
1053 *Representations*.